

# Guide for qjanno v1.0.0.0

## Contents

<b>1</b>	<b>Background</b>	<b>1</b>
<b>2</b>	<b>How does this work?</b>	<b>1</b>
<b>3</b>	<b>Installation</b>	<b>2</b>
<b>4</b>	<b>The CLI interface</b>	<b>2</b>
4.1	A basic example . . . . .	3
4.2	The .janno-crawling pseudo-functions . . . . .	3
4.3	CLI details . . . . .	4
4.4	The <code>-c</code> / <code>--showColumns</code> option . . . . .	5
<b>5</b>	<b>Query examples</b>	<b>5</b>

## 1 Background

qjanno started as a fork of the `qhs` software tool, which was, in turn, inspired by the CLI tool `q`. All of them enable SQL queries on delimiter-separated text files (e.g. `.csv` or `.tsv`). For qjanno we copied the source code of `qhs` v0.3.3 (MIT-License) and adjusted it to provide a smooth experience with a special kind of `.tsv` file: The Poseidon `.janno` file.

Unlike `trident` or `xerxes` qjanno does not have a complete understanding of the `.janno`-file structure, and (mostly) treats it like a normal `.tsv` file. It does not validate the files upon reading and takes them at face value. Still `.janno` files are given special consideration: With a set of pseudo-functions in the `FROM` field of the SQL query they can be searched recursively and loaded together into one table.

qjanno still supports most features of `qhs`, so it can still read `.csv` and `.tsv` files independently or in conjunction with `.janno` files (e.g. for `JOIN` operations).

## 2 How does this work?

On startup, qjanno creates an `SQLite` ([1]) database `in memory`. It then reads the requested, structured text files, attributes each column a type (either character or numeric) and writes the contents of the files to tables in the in-memory database. It finally sends the user-provided SQL query to the database, waits for the result, parses it and returns it on the command line.

The query gets pre-parsed to extract file names and then forwarded to an `SQLite` database server via the Haskell

library `sqlite-simple`. That means `qjanno` can parse and understand basic SQLite3 syntax, though not everything. `PRAGMA` functions, for example, are not available. The examples below show some of the available syntax, but they are not exhaustive. Trial and error is recommended to see what does and what does not work. Please report missing functionality in our [issue board on GitHub](#).

## 3 Installation

See the Poseidon website (<https://www.poseidon-adna.org/#/qjanno>) or the GitHub repository (<https://github.com/poseidon-framework/qjanno>) for up-to-date installation instructions.

## 4 The CLI interface

This is the CLI interface of `qjanno`:

```
Usage: qjanno [--version] [QUERY] [-q|--queryFile FILE] [-c|--showColumns]
        [-t|--tabSep] [--sep DELIM] [--noHeader] [--raw] [--noOutHeader]
```

Command line tool to allow SQL queries on `.janno` (and arbitrary `.csv` and `.tsv`) files.

Available options:

<code>-h,--help</code>	Show this help text
<code>--version</code>	Show <code>qjanno</code> version
<code>QUERY</code>	SQLite syntax query with paths to files for table names. See the online documentation for examples. The special table name syntax <code>'d(path1,path2,...)'</code> treats the paths ( <code>path1</code> , <code>path2</code> , ...) as base directories where <code>.janno</code> files are searched recursively. All detected <code>.janno</code> files are merged into one table and can thus be subjected to arbitrary queries.
<code>-q,--queryFile FILE</code>	Read query from the provided file.
<code>-c,--showColumns</code>	Don't run the query, but show all available columns in the input files.
<code>-t,--tabSep</code>	Short for <code>--sep '\$\t'</code> .
<code>--sep DELIM</code>	Input file field delimiter. Will be automatically detected if it's not specified.
<code>--noHeader</code>	Does the input file have no column names? They will be filled automatically with placeholders of the form <code>c1,c2,c3,...</code> .
<code>--raw</code>	Return the output table as <code>tsv</code> .
<code>--noOutHeader</code>	Remove the header line from the output.

This help can be accessed with `qjanno -h`. Running `qjanno` without any parameters does not work: The `QUERY` parameter is mandatory and the tool will fail with `Query cannot be empty`.

## 4.1 A basic example

A basic, working qjanno query could look like this:

```
$ qjanno "SELECT package_title,Poseidon_ID,Country \
        FROM d(2010_RasmussenNature,2012_MeyerScience)"
.----- .----- .-----
| package_title | Poseidon_ID | Country |
:===== :===== :=====
| 2010_RasmussenNature | Inuk.SG | Greenland |
| 2012_MeyerScience | A_Mbuti-5.DG | Congo |
| 2012_MeyerScience | A_Yoruba-4.DG | Nigeria |
| 2012_MeyerScience | A_Sardinian-4.DG | Italy |
| 2012_MeyerScience | A_French-4.DG | France |
| 2012_MeyerScience | A_Dinka-4.DG | Sudan |
| 2012_MeyerScience | A_Ju_hoan_North-5.DG | Namibia |
'-----'-----'-----'
```

Running qjanno with this query triggers the following process:

1. With `d(...)` in the `FROM` field, qjanno searches recursively for package-defining `POSEIDON.yml` files in the given base directories `2010_RasmussenNature` and `2012_MeyerScience`.
2. It finds the `.yml` files and reads some of their fields, including the `title`, the `packageVersion` and the `jannoFile` path. It then selects the latest version of each package.
3. With the relevant `.janno` file paths available, qjanno reads them, appends the `package_title`, `package_version` and `source_file` columns, merges them (simple row-bind), and orders their columns.
4. It then writes the resulting `.janno` table to the SQLite database in memory.
5. Now the actual query gets executed. In this case the `SELECT` statement includes three variables (column names): `package_title`, `Poseidon_ID` and `Country`. The database server returns these three columns for the merged `.janno` table.
6. qjanno finally prints the result in a neat, human readable format to the standard output.

## 4.2 The .janno-crawling pseudo-functions

`d(...)` is one of four mechanisms to search and load `.janno` files in the `FROM` field of the query:

- `d(<path_to_directory1>,<path_to_directory2>,...)`: With `d()`, qjanno (recursively) searches all package-defining `POSEIDON.yml` files in all listed directories and reads them to determine the latest package version. It then reads the `.janno` files associated with these latest package versions.
- `da(<path_to_directory1>,<path_to_directory2>,...)`: `da()` behaves just as `d()`, but it does not filter for the latest package version: It loads all packaged `.janno` files.
- `j(<path_to_directory1>,<path_to_directory2>,...)`: `j()` simply searches for files with the extension `.janno` in all listed directories and loads them regardless of whether they are part of a Poseidon package or not.
- `<path_to_one_janno_file>.janno`: Specific `.janno` files can be listed individually. They are identified as such by their `.janno` extension.

Multiple of these methods can be combined as a comma-separated list. Each respective mechanism then yields a

list of .janno file paths, and the list of lists is flattened to a simple list of paths. `qjanno` then reads all files in this combined list, merges them and makes them available for querying in the in-memory SQLite database.

!> Note that `FROM` field should not include any spaces – even in a comma-separated list. `qjanno` parses the `QUERY` using space as a separator.

### 4.3 CLI details

`qjanno` can not just read .janno files, but also arbitrary .csv and .tsv files. This option is triggered by providing file names (relative paths) in the `FROM` field of the query, not `d(...)`.

```
$ echo -e "Col1,Col2\nVal1,Val2\nVal3,Val4\n" > test.csv
$ qjanno "SELECT * FROM test.csv"
.------.------.-----
| source_file | Col1 | Col2 |
:=====:=====:=====:
| test.csv    | Val1 | Val2 |
| test.csv    | Val3 | Val4 |
'-----'-----'-----'
```

With these non-janno files `qjanno` automatically tries to detect the relevant separator. With `--sep` a delimiter can be specified explicitly, and the shortcut `-t` sets `--sep '$\t'` for tab-separated files.

```
$ echo -e "Col1\tCol2\nVal1\tVal2\nVal3\tVal4\n" > test.csv
$ qjanno "SELECT * FROM test.csv" -t # -t is optional
.------.------.-----
| source_file | Col1 | Col2 |
:=====:=====:=====:
| test_tab.csv | Val1 | Val2 |
| test_tab.csv | Val3 | Val4 |
'-----'-----'-----'
```

The `--noHeader` option allows to read files without headers, so column names. The columns are then automatically named `c1, c2, ... cN`:

```
$ echo -e "Val1,Val2\nVal3,Val4\n" > test.csv
$ qjanno "SELECT c1,c2 FROM test.csv" --noHeader
.------.-----
| c1 | c2 |
:=====:=====:
| Val1 | Val2 |
| Val3 | Val4 |
'-----'-----'
```

The remaining options concern the output: `--raw` returns the output table not in the neat, human-readable ASCII table layout, but in a simple .tsv format. `--noOutHeader` omits the header line in the output.

```
$ echo -e "Col1,Col2\nVal1,Val2\nVal3,Val4\n" > test.csv
$ qjanno "SELECT * FROM test.csv" --raw --noOutHeader
test.csv Val1 Val2
test.csv Val3 Val4
```

150 Note that these output options allow to directly prepare individual lists in trident's forgeScript selection language  
151 format:

```
152 $ qjanno "SELECT '<||Poseidon_ID||>' FROM d(2012_MeyerScience)" --raw --noOutHeader
153 <A_Mbuti-5.DG>
154 <A_Yoruba-4.DG>
155 <A_Sardinian-4.DG>
156 <A_French-4.DG>
157 <A_Dinka-4.DG>
158 <A_Ju_hoan_North-5.DG>
```

## 159 4.4 The `-c / --showColumns` option

160 `-c / --showColumns` is a special option that, when activated, makes qjanno return not the result of a given  
161 query, but an overview table with the columns available in all loaded tables/files for said query. That is helpful  
162 to get an overview what could actually be queried.

```
163 $ echo -e "Col1,Col2\nVal1,Val2\nVal3,Val4\n" > test.csv
164 $ qjanno "SELECT * FROM test.csv" -c
165 .------.------.------.
166 | Column | Path | qjanno Table name |
167 :=====:=====:=====:
168 | source_file | test.csv | test |
169 | Col1 | test.csv | test |
170 | Col2 | test.csv | test |
171 '-----'-----'-----'
```

172 This summary also includes the artificial, structurally cleaned table names assigned by `qjanno` before writing  
173 to the SQLite database. Often we can not simply use the file names as table names, because SQLite has strict  
174 naming requirements. File names or relative paths are generally invalid as table names and need to be replaced  
175 with a tidy string. These artificially generated names are mostly irrelevant from a user perspective – except a  
176 query involves multiple files, e.g. in a `JOIN` operation. See below for an example.

## 177 5 Query examples

178 The following examples show some of the functionality of the SQLite query language available through qjanno.  
179 See the [SQLite syntax documentation](#) for more details. They were prepared and tested in a clone of the Poseidon  
180 community archive.

### 181 Sub-setting with `WHERE`

182 Get all individuals (rows) in two Poseidon packages where UDG is set to 'minus'.

```
183 $ qjanno " \
184 SELECT package_title,Poseidon_ID,UDG \
185 FROM d(2010_RasmussenNature,2012_MeyerScience) \
186 WHERE UDG = 'minus' \
187 "
188 .------.-----.
```

```

189 | Poseidon_ID | UDG |
190 :=====:=====:
191 | Inuk.SG      | minus |
192 '-----'-----'

193 Get all individuals where Genetic_Sex is not 'F' and Country is 'Sudan'.

194 $ qjanno " \
195 SELECT Poseidon_ID,Country \
196 FROM d(2010_RasmussenNature,2012_MeyerScience) \
197 WHERE Genetic_Sex <> 'F' AND Country = 'Sudan' \
198 "
199 .-----.-----.
200 | Poseidon_ID | Country |
201 :=====:=====:
202 | A_Dinka-4.DG | Sudan |
203 '-----'-----'

204 Get all individuals where the the UDG column is not NULL or the Country is 'Sudan'.

205 $ qjanno " \
206 SELECT Poseidon_ID,Country \
207 FROM d(2010_RasmussenNature,2012_MeyerScience) \
208 WHERE UDG IS NOT NULL OR Country = 'Sudan' \
209 "
210 .-----.-----.
211 | Poseidon_ID | Country |
212 :=====:=====:
213 | Inuk.SG      | Greenland |
214 | A_Dinka-4.DG | Sudan |
215 '-----'-----'

216 Get all individuals where Nr_SNPs is equal to or bigger than 600,000.

217 $ qjanno " \
218 SELECT Poseidon_ID,Nr_SNPs \
219 FROM d(2010_RasmussenNature,2012_MeyerScience) \
220 WHERE Nr_SNPs >= 600000 \
221 "
222 .-----.-----.
223 | Poseidon_ID | Nr_SNPs |
224 :=====:=====:
225 | Inuk.SG      | 1101700 |
226 '-----'-----'

227 Ordering with ORDER BY

228 Order all individuals by Nr_SNPs.

229 $ qjanno " \
230 SELECT Poseidon_ID,Nr_SNPs \

```

```

231 FROM d(2010_RasmussenNature,2012_MeyerScience) \
232 ORDER BY Nr_SNPs \
233 "
234 .----- .----- .
235 |      Poseidon_ID      | Nr_SNPs |
236 :===== :===== :
237 | A_French-4.DG        | 592535 |
238 | A_Ju_hoan_North-5.DG | 593045 |
239 | A_Mbuti-5.DG         | 593057 |
240 | A_Dinka-4.DG         | 593076 |
241 | A_Yoruba-4.DG        | 593097 |
242 | A_Sardinian-4.DG     | 593109 |
243 | Inuk.SG              | 1101700 |
244 '-----'-----'

```

245 Order all individuals by Date\_BC\_AD\_Median in a descending ( **DESC** ) order. Date\_BC\_AD\_Median includes  
246 NULL values.

```

247 $ qjanno " \
248 SELECT Poseidon_ID,Date_BC_AD_Median \
249 FROM d(2010_RasmussenNature,2012_MeyerScience) \
250 ORDER BY Date_BC_AD_Median DESC \
251 "
252 .----- .----- .
253 |      Poseidon_ID      | Date_BC_AD_Median |
254 :===== :===== :
255 | Inuk.SG              | -1935             |
256 | A_Sardinian-4.DG     |                   |
257 | A_Yoruba-4.DG        |                   |
258 | A_Dinka-4.DG         |                   |
259 | A_Mbuti-5.DG         |                   |
260 | A_Ju_hoan_North-5.DG |                   |
261 | A_French-4.DG        |                   |
262 '-----'-----'

```

## 263 Reducing the number of return values with **LIMIT**

264 Only return the first three result individuals.

```

265 $ qjanno " \
266 SELECT Poseidon_ID,Group_Name \
267 FROM d(2010_RasmussenNature,2012_MeyerScience) \
268 LIMIT 3 \
269 "
270 .----- .----- .
271 | Poseidon_ID |      Group_Name      |
272 :===== :===== :
273 | Inuk.SG     | Greenland_Saqqaq.SG |

```

```

274 | A_Mbuti-5.DG | Ignore_Mbuti(discovery).DG |
275 | A_Yoruba-4.DG | Ignore_Yoruba(discovery).DG |
276 '-----'-----'

```

## 277 Combining tables with JOIN

278 For **JOIN** operations, SQLite requires table names to specify which columns are meant when combining multiple  
 279 tables with overlapping column names. See the option **-c / --showColumns** to get the relevant table names as  
 280 generated from the input paths.

```

281 $ echo -e "Poseidon_ID,MoreInfo\nInuk.SG,5\nA_French-4.DG,3\n" > test.csv

```

```

282
283 $ qjanno "SELECT * FROM d(2010_RasmussenNature,2012_MeyerScience)" -c
284 .-----'.-----'.-----'.
285 |          Column          |          Path          |
286 :=====:=====:=====:
287 | package_title            | d(2010_RasmussenNature,2012_MeyerScience) |
288 | package_version          | d(2010_RasmussenNature,2012_MeyerScience) |
289 | source_file              | d(2010_RasmussenNature,2012_MeyerScience) |
290 | Poseidon_ID              | d(2010_RasmussenNature,2012_MeyerScience) |
291 ...
292 -----'.
293 |          qjanno Table name          |
294 =====:
295 | d2010RasmussenNature2012MeyerScience |
296 | d2010RasmussenNature2012MeyerScience |
297 | d2010RasmussenNature2012MeyerScience |
298 | d2010RasmussenNature2012MeyerScience |
299 ...

```

```

300
301 $ qjanno "SELECT * FROM test.csv" -c
302 .-----'.-----'.-----'.
303 |   Column   |   Path   | qjanno Table name |
304 :=====:=====:=====:
305 | source_file | test.csv | test              |
306 | Poseidon_ID | test.csv | test              |
307 ...

```

308 Join the .janno files with the information in the test.csv file (by the **Poseidon\_ID** column).

```

309 $ qjanno " \
310 SELECT d2010RasmussenNature2012MeyerScience.Poseidon_ID,Country,MoreInfo \
311 FROM d(2010_RasmussenNature,2012_MeyerScience) \
312 INNER JOIN test.csv \
313 ON d2010RasmussenNature2012MeyerScience.Poseidon_ID = test.Poseidon_ID \
314 "
315 .-----'.-----'.-----'.
316 | Poseidon_ID | Country | MoreInfo |

```



```

317 :=====;=====;=====;
318 | Inuk.SG      | Greenland | 5      |
319 | A_French-4.DG | France   | 3      |
320 '-----'-----'-----'

```

## 321 Grouping data and applying aggregate functions

322 SQLite provides a number of aggregation functions: `avg(X)`, `count(*)`, `count(X)`, `group_concat(X)`,  
 323 `group_concat(X,Y)`, `max(X)`, `min(X)`, `sum(X)`. See the documentation [here](#). These functions can be well  
 324 combined with the `GROUP BY` operation.

325 Determine the minimal number of SNPs across all individuals.

```

326 $ qjanno "SELECT min(Nr_SNPs) AS n FROM d(2010_RasmussenNature,2012_MeyerScience)"
327 .-----
328 |   n   |
329 :=====;
330 | 592535 |
331 '-----'

```

332 Count the number of individuals per `Date_Type` group and calculate the average `Nr_SNPs` for both groups.

```

333 $ qjanno " \
334 SELECT Date_Type,count(*),avg(Nr_SNPs) \
335 FROM d(2010_RasmussenNature,2012_MeyerScience) \
336 GROUP BY Date_Type \
337 "
338 .----- .----- .-----
339 | Date_Type | count(*) | avg(Nr_SNPs) |
340 :=====;=====;=====;
341 | C14       | 1        | 1101700.0     |
342 | modern    | 6        | 592986.5      |
343 '-----'-----'-----'

```

344

- 345 [1] K. P. Gaffney, M. Prammer, L. Brasfield, D. R. Hipp, D. Kennedy, and J. M. Patel, "SQLite: Past, present, and future," *Proceedings of the VLDB Endowment*, vol. 15, no. 12, pp. 3535–3547, Aug. 2022, doi: [10.14778/3554821.3554842](https://doi.org/10.14778/3554821.3554842).