

Guide for trident v1.4.1.0

Contents

1	Installation	1
2	The trident CLI	2
2.1	General notes	4
2.1.1	Logging and command line output	4
2.1.2	Package duplicates and versions	4
2.1.3	Individual/sample duplicates	4
2.1.4	Group names in .fam files	5
2.1.5	Whitespaces in the .janno file	5
3	Package creation and manipulation commands	5
3.1	Init command	5
3.2	Fetch command	6
3.3	Forge command	8
3.3.1	The forge selection language	11
3.3.2	Treatment of the genotype data while merging	12
3.3.3	Treatment of the .janno file while merging	13
3.3.4	Treatment of the .ssf file while merging	14
3.3.5	Treatment of the .bib file while merging	14
3.3.6	Other options	14
3.4	Genoconvert command	15
3.5	Jannocoalesce command	16
3.6	Rectify command	17
4	Inspection commands	19
4.1	List command	19
4.2	Summarise command	20
4.3	Survey command	21
4.4	Validate command	21

1 Installation

See the Poseidon website (<https://www.poseidon-adna.org/#/trident>) or the GitHub repository (<https://github.com/poseidon-framework/poseidon-hs>) for up-to-date installation instructions.

33 2 The trident CLI

34 Trident is a command line software tool structured in multiple subcommands. If you installed it properly you
35 can call it on the command line by typing `trident`. This will show an overview of the general options and all
36 subcommands, which are explained in detail below.

```
Usage: trident [--version] [--logMode MODE | --debug] [--errLength INT]
        [--inPlinkPopName MODE] (COMMAND | COMMAND)
```

trident is a management and analysis tool for Poseidon packages. Report issues
here: <https://github.com/poseidon-framework/poseidon-hs/issues>

Available options:

<code>-h, --help</code>	Show this help text
<code>--version</code>	Show version number
<code>--logMode MODE</code>	How information should be reported: NoLog, SimpleLog, DefaultLog, ServerLog or VerboseLog. (default: DefaultLog)
<code>--debug</code>	Short for <code>--logMode VerboseLog</code> .
<code>--errLength INT</code>	After how many characters should a potential error message be truncated. "Inf" for no truncation. (default: CharCount 1500)
<code>--inPlinkPopName MODE</code>	Where to read the population/group name from the FAM file in Plink-format. Three options are possible: asFamily (default) asPhenotype asBoth.

Package creation and manipulation commands:

<code>init</code>	Create a new Poseidon package from genotype data
<code>fetch</code>	Download data from a remote Poseidon repository
<code>forge</code>	Select packages, groups or individuals and create a new Poseidon package from them
<code>genoconvert</code>	Convert the genotype data in a Poseidon package to a different file format
<code>jannocoalesce</code>	Coalesce information from one or multiple janno files to another one
<code>rectify</code>	Adjust POSEIDON.yml files automatically to package changes

Inspection commands:

<code>list</code>	List packages, groups or individuals from local or remote Poseidon repositories
<code>summarise</code>	Get an overview over the content of one or multiple Poseidon packages
<code>survey</code>	Survey the degree of context information completeness for Poseidon packages
<code>validate</code>	Check Poseidon packages or package components for

Trident allows to work directly with genotype data (see `-p` below), but its optimized for the interaction with Poseidon packages, which wrap and contextualize the data. Most trident subcommands therefore have a central parameter, called `--baseDir` or simply `-d` to specify one or more base directories to look for packages. For example, if all Poseidon packages live inside a repository at `/path/to/poseidon/packages` you would simply say `trident <subcommand> -d /path/to/poseidon/dirs/` and `trident` would automatically search all subdirectories inside of the repository for valid Poseidon packages (as identified by valid `POSEIDON.yml` files).

You can arrange a Poseidon repository in a hierarchical way. For example:

```

/path/to/poseidon/packages
  /modern
    /2019_poseidon_package1
    /2019_poseidon_package2
  /ancient
    /...
    /...
  /Reference_Genomes
    /...
    /...
```

You can use this structure to select only the level of packages you're interested in, even individual ones, and you can make use of the fact that `-d` can be given multiple times.

Being able to specify one or multiple repositories is often not enough, as you may have your own data to co-analyse with the main repository. This is easy to do, as you simply need to provide your own genotype data as yet another Poseidon package to be added to your `trident` command. For example, let's say you have genotype data in `EIGENSTRAT` format (`trident` supports `EIGENSTRAT` and `PLINK` as formats.):

```

~/my_project/my_project.geno
~/my_project/my_project.snp
~/my_project/my_project.ind
```

Then you can make that to a skeleton Poseidon package with the `init` command. You can also do it manually by simply adding a `POSEIDON.yml` file, with for example the following content:

```

poseidonVersion: 2.7.1
title: My_awesome_project
description: Unpublished genetic data from my awesome project
contributor:
  - name: Stephan Schiffels
    email: schiffels@institute.org
packageVersion: 0.1.0
lastModified: 2020-10-07
genotypeData:
  format: EIGENSTRAT
  genoFile: my_project.geno
  snpFile: my_project.snp
  indFile: my_project.ind
```

78 jannoFile: my_project.janno

79 bibFile: sources.bib

80 Two remarks: 1) all file paths are considered *relative* to the directory in which POSEIDON.yml resides. For this
81 example we assume that this file is added into the same directory as the three genotype files. 2) Besides the
82 genotype data files there are two (technically optional) files referenced by this example POSEIDON.yml file:
83 sources.bib and my_project.janno. Of course you can add them manually - `init` automatically creates empty
84 dummy versions.

85 Once you have set up your own Poseidon package (which is really only a skeleton so far), you can add it to your
86 trident analysis, by simply adding your project directory to the command using `-d`, for example:

87 trident list -d /path/to/poseidon/packages/modern \

88 -d /path/to/poseidon/packages/ReferenceGenomes

89 -d ~/my_project --packages

90 2.1 General notes

91 2.1.1 Logging and command line output

92 For all subcommands the general argument `--logMode` defines how trident reports messages (to stderr) on the
93 command line:

- 94 • *NoLog*: Hides all messages.
- 95 • *SimpleLog*: Plain and simple output to stderr.
- 96 • *DefaultLog*: Adds severity indicators before each message. (default setting)
- 97 • *ServerLog*: Additionally adds timestamps before each message.
- 98 • *VerboseLog*: Shows not just messages on the log levels **Info**, **Warning** and **Error** like the other modes, but
99 also on the more verbose level **Debug**. Use this for debugging.

100 `--debug` is short for `--logMode VerboseLog` to activate this important log level more easily.

101 2.1.2 Package duplicates and versions

- 102 • For **trident** multiple packages in a set of base directories can share the same **title**, if they have different
103 **packageVersion** numbers. If the version numbers are identical or missing, then **trident** stops with an
104 exception.
- 105 • The **trident** subcommands **genoconvert**, **list**, **rectify**, **survey** and **validate** by default consider all
106 versions of each Poseidon package in the given base directories. The `--onlyLatest` flag causes them to
107 instead only consider the latest versions.
- 108 • **fetch** and **forge** generally consider all package versions and their selection language (see below) allows
109 for detailed version handling.
- 110 • **summarize** and **jannocoalesce** always only consider the latest package versions.

111 2.1.3 Individual/sample duplicates

- 112 • Individual/sample names (**Poseidon_IDs**) within one package have to be unique, or trident will stop.
- 113 • We also discourage sample duplicates across packages in package repositories, but trident will generally
114 continue with them. **validate** will fail though, if the `--ignoreDuplicates` flag is not set.
- 115 • **forge** offers a special mechanism to resolve sample duplicates within its selection language.

116 2.1.4 Group names in .fam files

117 The .fam file of Plink-formatted genotype data is used inconsistently across different popular aDNA software
118 tools to store group/population name information. The (global) option `--inPlinkPopName` with the arguments
119 `asFamily` (default), `asPhenotype` and `asBoth` allows to control the reading of the population name from Plink
120 .fam files. The subcommands that write genotype data (`forge`, `genoconvert`) have a corresponding option
121 `--outPlinkPopName` to specify this for the output.

122 2.1.5 Whitespaces in the .janno file

123 While reading the .janno file `trident` trims all leading and trailing whitespaces around individual cells. Also
124 all instances of the `No-Break Space` unicode character will be removed. This means these whitespaces will not
125 be preserved when a package is `forged`.

126 3 Package creation and manipulation commands

127 3.1 Init command

128 `init` creates a new, valid Poseidon package from genotype data files. It adds a valid `POSEIDON.yml` file, a dummy
129 .janno file for context information and an empty .bib file for literature references.

130 Command line details

```
131 Usage: trident init ((-p|--genoOne FILE) | --inFormat FORMAT --genoFile FILE
132                  --snpFile FILE --indFile FILE) [--snpSet SET]
133                  (-o|--outPackagePath DIR) [-n|--outPackageName STRING]
134                  [--minimal]
```

135
136 Create a new Poseidon package from genotype data

137
138 Available options:

139	<code>-h,--help</code>	Show this help text
140	<code>-p,--genoOne FILE</code>	One of the input genotype data files. Expects .bed, 141 .bim or .fam for PLINK and .geno, .snp or .ind for 142 EIGENSTRAT. The other files must be in the same 143 directory and must have the same base name.
144	<code>--inFormat FORMAT</code>	The format of the input genotype data: EIGENSTRAT or 145 PLINK. Only necessary for data input with <code>--genoFile</code> 146 + <code>--snpFile</code> + <code>--indFile</code> .
147	<code>--genoFile FILE</code>	Path to the input geno file.
148	<code>--snpFile FILE</code>	Path to the input snp file.
149	<code>--indFile FILE</code>	Path to the input ind file.
150	<code>--snpSet SET</code>	The snpSet of the package: 1240K, HumanOrigins or 151 Other. Only relevant for data input with <code>-p --genoOne</code> 152 or <code>--genoFile</code> + <code>--snpFile</code> + <code>--indFile</code> , because the 153 packages in a <code>-d --baseDir</code> already have this 154 information in their respective <code>POSEIDON.yml</code> files. 155 (default: Other)

```

156 -o,--outPackagePath DIR Path to the output package directory.
157 -n,--outPackageName STRING
158                               The output package name. This is optional: If no name
159                               is provided, then the package name defaults to the
160                               basename of the (mandatory) --outPackagePath
161                               argument. (default: Nothing)
162 --minimal                     Should the output data be reduced to a necessary
163                               minimum and omit empty scaffolding?
164
164 The command
165
165 trident init \
166   --inFormat EIGENSTRAT/PLINK \
167   --genoFile path/to/geno_file \
168   --snpFile path/to/snp_file \
169   --indFile path/to/ind_file \
170   --snpSet 1240K|HumanOrigins|Other \
171   -o path/to/new_package_name
172
172 requires the format (--inFormat) of your input data (either EIGENSTRAT or PLINK), the paths to the respective
173 files (--genoFile, --snpFile, --indFile), and optionally the “shape” of these files (--snpSet), so if they cover
174 the 1240K, the HumanOrigins or an Other SNP set. A simpler interface is available with -p (+ --snpSet).

```

	EIGENSTRAT	PLINK
genoFile	.geno	.bed
snpFile	.snp	.bim
indFile	.ind	.fam

```

175 The output package of init is created as a new directory -o, which should not already exist, and gets the
176 package title corresponding to the basename of -o. You can also set the title explicitly with -n. The --minimal
177 flag causes init to create a minimal package with a very basic POSEIDON.yml and no .bib and .janno files.

```

178 3.2 Fetch command

```

179 fetch allows to download Poseidon packages from a remote Poseidon server via a Web API. This server provides
180 all packages in the Poseidon public archives.

```

```

181 Command line details

```

```

182 Usage: trident fetch (-d|--baseDir DIR)
183         (--downloadAll |
184         (--fetchFile FILE | (-f|--fetchString DSL)))
185         [--remoteURL URL] [--archive STRING]

```

```

186
187 Download data from a remote Poseidon repository

```

```

188
189 Available options:

```

```

190 -h,--help           Show this help text

```

```

191  -d,--baseDir DIR          A base directory to search for Poseidon packages.
192  --downloadAll             Download all packages the server is offering.
193  --fetchFile FILE         A file with a list of packages. Works just as -f, but
194                           multiple values can also be separated by newline, not
195                           just by comma. -f and --fetchFile can be combined.
196  -f,--fetchString DSL     List of packages to be downloaded from the remote
197                           server. Package names should be wrapped in asterisks:
198                           *package_title*. You can combine multiple values with
199                           comma, so for example: "*package_1*, *package_2*,
200                           *package_3*". fetchString uses the same parser as
201                           forgeString, but does not allow excludes. If groups
202                           or individuals are specified, then packages which
203                           include these groups or individuals are included in
204                           the download.
205  --remoteURL URL          URL of the remote Poseidon server.
206                           (default: "https://server.poseidon-adna.org")
207  --archive STRING         The name of the Poseidon package archive that should
208                           be queried. If not given, then the query falls back
209                           to the default archive of the server selected with
210                           --remoteURL. See the archive documentation at
211                           https://www.poseidon-adna.org/#/archive_overview for
212                           a list of archives currently available from the
213                           official Poseidon Web API. (default: Nothing)
214
215  It works with
216
217  trident fetch -d ... -d ... \
218    -f "*package_title_1*,*package_title_2-1.0.1*,group_name,<individual1>"
219
220  and the entities you want to download must be listed either in a simple string of comma-separated values, which
221  can be passed via -f/--fetchString, or in a text file (--fetchFile). Entities are then combined from these
222  sources.
223
224  Entities are specified using a special syntax (see also the documentation of forge below): packages are wrapped
225  in asterisks, with or without version appended after a dash (e.g. *package_title* or *package_title-1.2.3*),
226  group names are spelled as is, and individual names are wrapped in angular brackets (e.g. <individual1>).
227  Fetch will figure out which packages need to be downloaded to include all specified entities. --downloadAll,
228  which can be given instead of -f and --fetchFile, causes fetch to download all packages from the server. The
229  downloaded packages are added in the first (!) -d directory (which gets created if it doesn't exist), but downloads
230  are only performed if the respective packages are not already present in the latest version in any of the -d dirs.
231
232  Note that trident fetch makes most sense in combination with trident list --remote: First one can inspect
233  what is available on the server, then one can create a custom fetch command.
234
235  fetch also has the optional arguments --remote https://... to name an alternative Poseidon server and
236  --archive to select a specific Poseidon public archive on the server.

```

231 3.3 Forge command

232 **forge** creates new Poseidon packages by extracting and merging packages, populations and individuals/samples
233 from your Poseidon repositories.

234 Command line details

```
235 Usage: trident forge ((-d|--baseDir DIR) |  
236                      ((-p|--genoOne FILE) | --inFormat FORMAT --genoFile FILE  
237                      --snpFile FILE --indFile FILE) [--snpSet SET])  
238                      [--forgeFile FILE | (-f|--forgeString DSL)]  
239                      [--selectSnps FILE] [--intersect] [--outFormat FORMAT]  
240                      [--minimal] [--onlyGeno] (-o|--outPackagePath DIR)  
241                      [-n|--outPackageName STRING] [--packagewise]  
242                      [--outPlinkPopName MODE]
```

244 Select packages, groups or individuals and create a new Poseidon package from
245 them

247 Available options:

248 -h,--help	Show this help text
249 -d,--baseDir DIR	A base directory to search for Poseidon packages.
250 -p,--genoOne FILE	One of the input genotype data files. Expects .bed, 251 .bim or .fam for PLINK and .geno, .snp or .ind for 252 EIGENSTRAT. The other files must be in the same 253 directory and must have the same base name.
254 --inFormat FORMAT	The format of the input genotype data: EIGENSTRAT or 255 PLINK. Only necessary for data input with --genoFile 256 + --snpFile + --indFile.
257 --genoFile FILE	Path to the input geno file.
258 --snpFile FILE	Path to the input snp file.
259 --indFile FILE	Path to the input ind file.
260 --snpSet SET	The snpSet of the package: 1240K, HumanOrigins or 261 Other. Only relevant for data input with -p --genoOne 262 or --genoFile + --snpFile + --indFile, because the 263 packages in a -d --baseDir already have this 264 information in their respective POSEIDON.yml files. 265 (default: Other)
266 --forgeFile FILE	A file with a list of packages, groups or individual 267 samples. Works just as -f, but multiple values can 268 also be separated by newline, not just by comma. 269 Empty lines are ignored and comments start with "#", 270 so everything after "#" is ignored in one line. 271 Multiple instances of -f and --forgeFile can be 272 given. They will be evaluated according to their 273 input order on the command line.
274 -f,--forgeString DSL	List of packages, groups or individual samples to be

combined in the output package. Packages follow the syntax `*package_title*`, populations/groups are simply `group_id` and individuals `<individual_id>`. You can combine multiple values with comma, so for example: `"*package_1*, <individual_1>, <individual_2>, group_1"`. Duplicates are treated as one entry. Negative selection is possible by prepending "-" to the entity you want to exclude (e.g. `"*package_1*, -<individual_1>, -group_1"`). `forge` will apply excludes and includes in order. If the first entity is negative, then `forge` will assume you want to merge all individuals in the packages found in the `baseDirs` (except the ones explicitly excluded) before the exclude entities are applied. An empty `forgeString` (and no `--forgeFile`) will therefore merge all available individuals. If there are individuals in your input packages with equal individual id, but different main group or source package, they can be specified with the special syntax `"<package:group:individual>"`.

`--selectSnps FILE` To extract specific SNPs during this `forge` operation, provide a Snp file. Can be either Eigenstrat (file ending must be `'.snp'`) or Plink (file ending must be `'.bim'`). When this option is set, the output package will have exactly the SNPs listed in this file. Any SNP not listed in the file will be excluded. If option `'--intersect'` is also set, only the SNPs overlapping between the SNP file and the forged packages are output. (default: Nothing)

`--intersect` Whether to output the intersection of the genotype files to be forged. The default (if this option is not set) is to output the union of all SNPs, with genotypes defined as missing in those packages which do not have a SNP that is present in another package. With this option set, the forged dataset will typically have fewer SNPs, but less missingness.

`--outFormat FORMAT` The format of the output genotype data: EIGENSTRAT or PLINK. (default: PLINK)

`--minimal` Should the output data be reduced to a necessary minimum and omit empty scaffolding?

`--onlyGeno` Should only the resulting genotype data be returned? This means the output will not be a Poseidon package.

`-o,--outPackagePath DIR` Path to the output package directory.

`-n,--outPackageName STRING` The output package name. This is optional: If no name

```

320         is provided, then the package name defaults to the
321         basename of the (mandatory) --outPackagePath
322         argument. (default: Nothing)
323     --packagewise      Skip the within-package selection step in forge. This
324                       will result in outputting all individuals in the
325                       relevant packages, and hence a superset of the
326                       requested individuals/groups. It may result in better
327                       performance in cases where one wants to forge entire
328                       packages or almost entire packages. Details: Forge
329                       conceptually performs two types of selection: First,
330                       it identifies which packages in the supplied base
331                       directories are relevant to the requested forge, i.e.
332                       whether they are either explicitly listed using
333                       *PackageName*, or because they contain selected
334                       individuals or groups. Second, within each relevant
335                       package, individuals which are not requested are
336                       removed. This option skips only the second step, but
337                       still performs the first.
338     --outPlinkPopName MODE Where to write the population/group name into the FAM
339                           file in Plink-format. Three options are possible:
340                           asFamily (default) | asPhenotype | asBoth. See also
341                           --inPlinkPopName.
342
343     forge can be used with
344
345     trident forge -d ... -d ... \
346         -f "*package_name*, group_id, <individual_id>" \
347         -o path/to/new_package_name
348
349     where the entities (packages, groups/populations, individuals/samples) you want in the output package can be
350     denoted either as a string on the command line (-f/--forgeString), or in an input text file (--forgeFile).
351     See the section below for the syntax of this selection language. Do not forget to wrap the --forgeString query
352     in quotes.
353
354     Including one or multiple Poseidon packages with -d is not the only way to include data for a forge operation.
355     It is also possible to consider unpackaged genotype data directly with -p (+ --snpSet) or --inFormat +
356     --genoFile + --snpFile + --indFile (+ --snpSet). This makes the following example possible, where we
357     merge data from one Poseidon package and two genotype datasets to get a new EIGENSTRAT dataset.
358
359     trident forge \
360         -d 2017_GonzalesFortesCurrentBiology \
361         -p 2018_VeeramahPNAS/2018_VeeramahPNAS.fam \
362         --inFormat PLINK \
363         --genoFile 2017_HaberAJHG/2017_HaberAJHG.bed \
364         --snpFile 2017_HaberAJHG/2017_HaberAJHG.bim \
365         --indFile 2017_HaberAJHG/2017_HaberAJHG.fam \
366         -f "<STR241.SG>,<ERS1790729.SG>,Iberia_HG.SG" \
367         -o testpackage \

```

```

363 --outFormat EIGENSTRAT \
364 --onlyGeno

```

365 3.3.1 The forge selection language

366 The text in `--forgeString`, `--forgeFile` (and with limited syntax also in `--fetchString` and `--fetchFile`)
367 are parsed as a domain specific query language that describes precisely which entities should be compiled in
368 the output package of a given **forge** operation. The language has multiple syntactic elements and a specific
369 evaluation logic.

370 In general a `--forgeString` query consists of multiple entities, separated by `,`. The main entities are Poseidon
371 packages, groups/populations and individuals/samples:

- 372 • Each package title is surrounded by `*`: `*package*`. That means if you want all individuals of the Poseidon
373 package `2019_Jeong_InnerEurasia` in the output package you would add `*2019_Jeong_InnerEurasia*`
374 to the query.
- 375 • Groups/populations are not specially marked: `group`. So to get all individuals of the group
376 `Swiss_Roman_period`, you would simply add `Swiss_Roman_period`.
- 377 • Individuals/samples are surrounded by `<` and `>`: `<individual>`. `ALA026` therefore becomes `<ALA026>`. A sec-
378 ond way to denote individuals is with the more verbose and specific syntax `<package:group:individual>`.
379 Such defined individuals take precedence over differently defined ones (so: directly with `<individual>` or
380 as a subset of `*package*` or `group`). This allows to resolve duplication issues precisely – at least in cases
381 where the duplicated individuals differ in source package or primary group.
- 382 • Package versions can be appended to package names, such as `*package-1.2.3*`.
- 383 • This also works with the verbose individual syntax: `<package-1.2.3:group:individual>`.

384 In the `--forgeFile` each line is treated as a separate `forgeString`, empty lines are ignored and `#` symbols start
385 comments. So this is a valid example of a `forgeFile`:

```

386 # Packages
387 *package1*, *package2-1.2.3*
388
389 # Groups and individuals from other packages beyond package1 and package2
390 group1, <individual1>, group2, <individual2>, <pac1:group2:individual3>
391
392 # group2 has two outlier individuals that should be ignored
393 -<individual1> # This one has very low coverage
394 -<pac2:group3:individual4> # This one is from a different time period

```

395 By prepending `-` to entities, we can exclude them from the forged package (this feature is not avail-
396 able for `fetch`). **forge** figures out the final list of samples to include by executing all **forge**-entities in
397 order. So an entity list `*PackageA*,-<Individual1>,GroupA` may result in a different outcome than
398 `*PackageA*,GroupA,-<Individual1>`, depending on whether `<Individual1>` belongs to `GroupA` or not.

399 If the **forge** entity list starts with a negative entity, or if the entity list is empty, **forge** will implicitly assume
400 you want to include all individuals in all **latest** versions of packages found in the base directories (except the
401 ones explicitly excluded, of course).

402 The specific semantics of the various ways to include or exclude entities are:

3.3.1.1 Inclusion queries

- ***Pac1***: Select all individuals in the latest version of package “Pac1”
- ***Pac1-1.0.1***: Select all individuals in package “Pac1” with version “1.0.1”
- **Group1**: Select all individuals associated with “Group1” in all latest versions of all packages
- **<Ind1>**: Select the individual named “Ind1”, searching in all latest packages.
- **<Pac1:Group1:Ind1>**: Select the individual named “Ind1” associated with “Group1” in the latest version of package “Pac1”
- **<Pac1-1.0.1:Group1:Ind1>**: Select the individual named “Ind1” associated with “Group1” in the package “Pac1” with version “1.0.1”

3.3.1.2 Exclusion queries

- **-*Pac1***: Remove all individuals in all versions of package “Pac1”
- **-*Pac1-1.0.1***: Remove only individuals in package “Pac1” with version “1.0.1” (but leave other versions in)
- **-Group1**: Remove all individuals associated with “Group1” in all versions of all packages (not just the latest)
- **-<Ind1>**: Remove all individuals named “Ind1” in all versions of all packages (not just the latest).
- **-<Pac1:Group1:Ind1>**: Remove the individual named “Ind1” associated with “Group1”, searching in all versions of package “Pac1”
- **-<Pac1-1.0.1:Group1:Ind1>**: Remove the individual named “Ind1” associated with “Group1”, but only if they are in “Pac1” with version “1.0.1”

If a query results in multiple individuals with the same name, forge will throw an error.

3.3.2 Treatment of the genotype data while merging

Forge performs a series of steps to merge the genotype data of multiple source files:

1. Genotype data from each package is streamed in parallel. Because our packages may have different SNP locations (specified by chromosome-position pairs) listed in their `.bim/.snp` file, we first perform a zipping-operation, whose behaviour depends on whether `--intersect` is set or not. Without `--intersect`, any SNP position listed in any package will be forwarded to the output, with missing values being filled in in all packages that do not list that particular SNP. With `--intersect`, only SNP positions that are present in all packages are considered. Note that relevant for this step is only whether a given SNP position is part of the genotype data, not whether the actual genotypes are missing or not.
2. At each SNP, the consensus alleles are selected, by collecting all reference and alternative alleles from all sources. If more than two non-dummy alleles (alleles different from N) are present in that collection, an error is thrown. If exactly two non-dummy alleles are present (which should be the case for binary SNPs), the two alleles are declared “reference” and “alternative” alleles for the output. If only one non-dummy allele is present, it is set to be the reference allele, and “N” is set to be the alternative.
3. All source genotype data is then read and recoded in terms of the two chosen consensus alleles. This will make sure that source data with flipped reference and alternative allele gets correctly merged in.
4. SNP IDs, as part of PLINK `.bim` files are checked across the source files. If all SNP IDs for a given SNP are missing, then the result will also be missing. If there is only one SNP ID present in some or all source packages, that ID gets forwarded to the output. In the (unusual) case that there are multiple different non-missing SNP ids (of the form “rs” followed by a number), then a debug warning is output (which gets

444 printed to the screen when `--logMode DEBUG` is selected), and simply the first value is chosen to be output
 445 into the forged `.bim` file. We decided not to throw an error in that case, because we consider the physical
 446 position of the SNP (specified by Chromosome and position) to be definitive, and the SNP ID to be of
 447 secondary importance.
 448 5. Genetic positions, as part of PLINK `.bim` files are checked in a similar manner, with “0.0” being interpreted
 449 as missing.

450 3.3.3 Treatment of the `.janno` file while merging

451 `forge` merges and subsets `.janno` files along with the genotype data. If a package lacks a `.janno` file, then a basic
 452 one will be created internally based on the information in the genotype data, and used for the output. Missing
 453 columns across packages will be filled with `n/a`.

454 For merging two `.janno` files **A** and **B** the following rules apply regarding undefined, arbitrary additional columns:

- 455 • If **A** has an additional column which is not in **B** then empty cells in the rows imported from **B** are filled
 456 with `n/a`.
- 457 • If **A** and **B** share additional columns with identical column name, then they are treated as semantically
 458 identical units and merged accordingly.
- 459 • In the resulting `.janno` file, all additional columns from both **A** and **B** are sorted alphabetically and
 460 appended after the normal, specified variables.

461 The following example illustrates the described behaviour:

462 A.janno

Poseidon_ID	Group_Name	Genetic_Sex	AdditionalColumn1	AdditionalColumn2
XXX011	POP1	M	A	D
XXX012	POP2	F	B	E
XXX013	POP1	M	C	F

463 B.janno

Poseidon_ID	Group_Name	Genetic_Sex	AdditionalColumn3	AdditionalColumn2
YYY022	POP5	F	G	J
YYY023	POP5	F	H	K
YYY024	POP5	M	I	L

464 A.janno + B.janno

Poseidon_ID	Group_Name	Genetic_Sex	AdditionalColumn1	AdditionalColumn2	AdditionalColumn3
XXX011	POP1	M	A	D	n/a
XXX012	POP2	F	B	E	n/a
XXX013	POP1	M	C	F	n/a
YYY022	POP5	F	n/a	J	G
YYY023	POP5	F	n/a	K	H

Poseidon_ID	Group_Name	Genetic_Sex	AdditionalColumn1	AdditionalColumn2	AdditionalColumn3
YYY024	POP5	M	n/a	L	I

3.3.4 Treatment of the .ssf file while merging

The Sequencing Source File (short .ssf file) is forged in exactly the same way as the janno file. SSF files that are present are included in the forge product in the way that the user expects, following selection of those entities which are listed in the `poseidon_IDs` columns of the SSF files. Columns that are only present in some packages, including those not defined by our [Schema] are also included in the forged product in the same way as described for Janno above.

3.3.5 Treatment of the .bib file while merging

In the forge process all relevant samples for the output package are determined. This includes their .janno entries and therefore the information on the publication keys documented for them in the .janno `Publication` column. The output .bib file compiles only the relevant references for the samples in the output package. It includes the references exactly once and is sorted alphabetically (by key).

3.3.6 Other options

Just as for `init` the output package of `forge` is created as a new directory `-o`. The title can also be explicitly defined with `-n`.

`--minimal` allows for the creation of a minimal output package without `.bib` and `.janno`. This is especially useful for data analysis pipelines, where only the genotype data is required. Even more basic output comes with `--onlyGeno`, which means that only the genotype data is returned without any Poseidon package.

`forge` has a an optional flag `--intersect`, that defines, if the genotype data from different packages should be merged with an **union** or an **intersect** operation. The default (if this option is not set) is to output the union of all SNPs, with genotypes defined as missing in samples from packages which do not have a SNP that is present in another package. With this option set, on the other hand, the forged dataset will typically have fewer SNPs, but less missingness.

`--intersect` also influences the automatic determination of the `snpSet` field in the `POSEIDON.yml` file for the resulting package. If the `snpSets` of all input packages are identical, then the resulting package will just inherit this configuration. Otherwise `forge` applies the following pairwise merging logic:

Input snpSet A	Input snpSet B	<code>--intersect</code>	Ouput snpSet
Other	*	*	Other
1240K	HumanOrigins	True	HumanOrigins
1240K	HumanOrigins	False	1240K

`--selectSnps` allows to provide `forge` with a SNP file in EIGENSTRAT (`.snp`) or PLINK (`.bim`) format to create a package with a specific selection. When this option is set, the output package will have exactly the SNPs listed in this file. Any SNP not listed in the file will be excluded. If `--intersect` is also set, only the SNPs overlapping between the SNP file and the forged packages are output.

494 Merging genotype data across different data sources and file formats is tricky. **forge** is more verbose about
495 potential issues, if the **--logMode** flag is set to **VerboseLog**.

496 The **--onlyGeno** command specifies that only genotype data should be output, not an entire Poseidon package.

497 With **--packagewise** the within-package selection step in **forge** can be skipped. This will result in outputting
498 all individuals in the relevant packages, and hence a superset of the requested individuals/groups. It may result
499 in better performance in cases where one wants to forge entire packages.

500 3.4 Genoconvert command

501 **genoconvert** converts the genotype data in a Poseidon package to a different file format. The respective entries
502 in the POSEIDON.yml file are changed accordingly.

503 Command line details

```
504 Usage: trident genoconvert ((-d|--baseDir DIR) |  
505                             ((-p|--genoOne FILE) | --inFormat FORMAT  
506                             --genoFile FILE --snpFile FILE --indFile FILE)  
507                             [--snpSet SET]) --outFormat FORMAT [--onlyGeno]  
508                             [-o|--outPackagePath DIR] [--removeOld]  
509                             [--outPlinkPopName MODE] [--onlyLatest]
```

511 Convert the genotype data in a Poseidon package to a different file format

513 Available options:

514 -h,--help	Show this help text
515 -d,--baseDir DIR	A base directory to search for Poseidon packages.
516 -p,--genoOne FILE	One of the input genotype data files. Expects .bed, 517 .bim or .fam for PLINK and .geno, .snp or .ind for 518 EIGENSTRAT. The other files must be in the same 519 directory and must have the same base name.
520 --inFormat FORMAT	The format of the input genotype data: EIGENSTRAT or 521 PLINK. Only necessary for data input with --genoFile 522 + --snpFile + --indFile.
523 --genoFile FILE	Path to the input geno file.
524 --snpFile FILE	Path to the input snp file.
525 --indFile FILE	Path to the input ind file.
526 --snpSet SET	The snpSet of the package: 1240K, HumanOrigins or 527 Other. Only relevant for data input with -p --genoOne 528 or --genoFile + --snpFile + --indFile, because the 529 packages in a -d --baseDir already have this 530 information in their respective POSEIDON.yml files. 531 (default: Other)
532 --outFormat FORMAT	the format of the output genotype data: EIGENSTRAT or 533 PLINK.
534 --onlyGeno	Should only the resulting genotype data be returned? 535 This means the output will not be a Poseidon package.

```

536 -o,--outPackagePath DIR Path to the output package directory. This is
537 optional: If no path is provided, then the output is
538 written to the directories where the input genotype
539 data file (.bed/.geno) is stored. (default: Nothing)
540 --removeOld Remove the old genotype files when creating the new
541 ones.
542 --outPlinkPopName MODE Where to write the population/group name into the FAM
543 file in Plink-format. Three options are possible:
544 asFamily (default) | asPhenotype | asBoth. See also
545 --inPlinkPopName.
546 --onlyLatest Consider only the latest versions of packages, or the
547 groups and individuals within the latest versions of
548 packages, respectively.

```

549 With the default setting

```

550 trident genoconvert -d ... -d ... --outFormat EIGENSTRAT|PLINK

```

551 all packages in `-d` will be converted to the desired `--outFormat` (either `EIGENSTRAT` or `PLINK`), if the data is
552 not already in this format. This includes updating the respective `POSEIDON.yml` files.

553 The “old” data is not deleted, but kept around. That means conversion can result in a package with both `PLINK`
554 and `EIGENSTRAT` data, but only one is linked in the `POSEIDON.yml` file, and that is what will be used by
555 `trident`. To delete the old data in the conversion you can add the `--removeOld` flag.

556 Instead of `-d` to change Poseidon packages, the `-p` (+ `--snpSet`) or `--inFormat` + `--genoFile` + `--snpFile`
557 + `--indFile` (+ `--snpSet`) allow to directly convert genotype data that is not wrapped in a Poseidon package
558 and store it to a directory given in `-o`. See this example:

```

559 trident genoconvert \
560   -p 2018_Mittnik_Baltic/Mittnik_Baltic.bed \
561   --outFormat EIGENSTRAT
562   -o my_directory

```

563 3.5 Jannocoalesce command

564 `jannocoalesce` merges information from one or multiple source `.janno` files into a target `.janno` file.

565 Command line details

```

566 Usage: trident jannocoalesce ((-s|--sourceFile FILE) | (-d|--baseDir DIR))
567                               (-t|--targetFile FILE) [-o|--outFile FILE]
568                               [--includeColumns ARG | --excludeColumns ARG]
569                               [-f|--force] [--sourceKey ARG] [--targetKey ARG]
570                               [--stripIdRegex ARG]

```

571

572 Coalesce information from one or multiple `janno` files to another one

573

574 Available options:

```

575 -h,--help          Show this help text
576 -s,--sourceFile FILE The source .janno file.

```



```

577 -d,--baseDir DIR      A base directory to search for Poseidon packages.
578 -t,--targetFile FILE  The target .janno file to fill.
579 -o,--outFile FILE     An optional file to write the results to. If not
580                       specified, change the target file in place.
581                       (default: Nothing)
582 --includeColumns ARG  A comma-separated list of .janno column names to
583                       coalesce. If not specified, all columns that can be
584                       found in the source and target will get filled.
585 --excludeColumns ARG  A comma-separated list of .janno column names NOT to
586                       coalesce. All columns that can be found in the source
587                       and target will get filled, except the ones listed
588                       here.
589 -f,--force            With this option, potential non-missing content in
590                       target columns gets overridden with non-missing
591                       content in source columns. By default, only missing
592                       data gets filled-in.
593 --sourceKey ARG       The .janno column to use as the source key.
594                       (default: "Poseidon_ID")
595 --targetKey ARG       The .janno column to use as the target key.
596                       (default: "Poseidon_ID")
597 --stripIdRegex ARG    An optional regular expression to identify parts of
598                       the IDs to strip before matching between source and
599                       target. Uses POSIX Extended regular expressions.
600
600 A most basic run may just include two arguments:
601
601 trident jannocoalesce \
602   --sourceFile path/to/source.janno \
603   --targetFile path/to/target.janno
604
604 jannocoalesce generally works by reading a source .janno file with -s|--sourceFile (or all .janno files in a
605 -d|--baseDir) and a target .janno file with -t|--targetFile.
606
606 It then merges these files by a key column, which can be selected with --sourceKey and --targetKey. The
607 default for both of these key columns is the Poseidon_ID. In case the entries in the key columns slightly and
608 systematically differ, e.g. because the Poseidon_IDs in either have a special suffix (for example _SG), then the
609 --stripIdRegex option allows to strip these with a regular expression to thus match the keys.
610
610 jannocoalesce generally attempts to fill all empty cells in the target .janno file with information from the
611 source. --includeColumns and --excludeColumns allow to select specific columns for which this should be
612 done. In some cases it may be desirable to not just fill empty fields in the target, but overwrite the information
613 already there with the -f|--force option. If the target file should be preserved, then the output can be directed
614 to a new output .janno file with -o|--outFile.

```

615 3.6 Rectify command

```

616 rectify automatically harmonizes POSEIDON.yml files of one or multiple packages. This is not an automatic
617 update from one Poseidon version to the next, but rather a clean-up wizard after manual modifications.

```

618 Command line details

```
619 Usage: trident rectify (-d|--baseDir DIR) [--ignorePoseidonVersion]
620                [--poseidonVersion ?.??.?]
621                [--packageVersion VPART [--logText STRING]]
622                [--checksumAll | [--checksumGeno] [--checksumJanno]
623                [--checksumSSF] [--checksumBib]]
624                [--newContributors DSL] [--onlyLatest]
```

625
626 Adjust POSEIDON.yml files automatically to package changes

627 Available options:

```
629  -h,--help           Show this help text
630  -d,--baseDir DIR    A base directory to search for Poseidon packages.
631  --ignorePoseidonVersion Read packages even if their poseidonVersion is not
632                      compatible with trident.
633  --poseidonVersion ?.??.? Poseidon version the packages should be updated to:
634                      e.g. "2.5.3".
635  --packageVersion VPART Part of the package version number in the
636                      POSEIDON.yml file that should be updated: Major,
637                      Minor or Patch (see https://semver.org).
638  --logText STRING    Log text for this version in the CHANGELOG file.
639  --checksumAll       Update all checksums.
640  --checksumGeno      Update genotype data checksums.
641  --checksumJanno     Update .janno file checksum.
642  --checksumSSF       Update .ssf file checksum
643  --checksumBib       Update .bib file checksum.
644  --newContributors DSL Contributors to add to the POSEIDON.yml file in the
645                      form "[Firstname Lastname](Email address);...".
646  --onlyLatest        Consider only the latest versions of packages, or the
647                      groups and individuals within the latest versions of
648                      packages, respectively.
```

649 It can be called with a lot of optional arguments. Note that `rectify` by default does **not** apply any changes if
650 none of these arguments are set.

```
651 trident rectify -d ... -d ... \
652   --poseidonVersion "X.X.X" \
653   --packageVersion Major|Minor|Patch \
654   --logText "short description of the update" \
655   --checksumAll \
656   --newContributors "[Firstname Lastname](Email address);..."
```

657 The following arguments determine which fields of the POSEIDON.yml file should be modified:

- 658 • `--poseidonVersion` allows a simple change of the `poseidonVersion` field in the POSEIDON.yml file.
- 659 • `--packageVersion` increments the package version number in the first, the second or the third position.
- 660 It can optionally be called with `--logText`, which appends an entry to the CHANGELOG file for the

661 respective package version update. `--logText` also creates a new CHANGELOG file if it does not exist
662 yet.

- 663 • `--checksumGeno`, `--checksumJanno`, `--checksumSSF` and `--checksumBib` add or modify the respective
664 checksum fields in the POSEIDON.yml file. `--checksumAll` is a wrapper to call all of them at once.
- 665 • `--newContributors` adds new contributors.

666 :warning: As `rectify` reads and rewrites POSEIDON.yml files, it may change their inner order, layout or even
667 content (e.g. if they have fields which are not in the POSEIDON.yml specification). Create a backup of the
668 POSEIDON.yml file before running `rectify` if you are uncertain if this might affect you negatively.

669 4 Inspection commands

670 4.1 List command

671 `list` lists packages, groups and individuals of the datasets you use, or of the packages available on the server.

672 Command line details

```
673 Usage: trident list ((-d|--baseDir DIR) | --remote [--remoteURL URL]
674                   [--archive STRING])
675                   (--packages | --groups | --individuals
676                   [-j|--jannoColumn COLNAME]) [--raw] [--onlyLatest]
```

678 List packages, groups or individuals from local or remote Poseidon
679 repositories

681 Available options:

682 <h1>-h,--help</h1>	Show this help text
683 <h1>-d,--baseDir DIR</h1>	A base directory to search for Poseidon packages.
684 <h1>--remote</h1>	List packages from a remote server instead the local 685 file system.
686 <h1>--remoteURL URL</h1>	URL of the remote Poseidon server. 687 (default: "https://server.poseidon-adna.org")
688 <h1>--archive STRING</h1>	The name of the Poseidon package archive that should 689 be queried. If not given, then the query falls back 690 to the default archive of the server selected with 691 --remoteURL. See the archive documentation at 692 https://www.poseidon-adna.org/#/archive_overview for 693 a list of archives currently available from the 694 official Poseidon Web API. (default: Nothing)
695 <h1>--packages</h1>	List all packages.
696 <h1>--groups</h1>	List all groups, ignoring any group names after the 697 first as specified in the .janno-file.
698 <h1>--individuals</h1>	List all individuals/samples.
699 <h1>-j,--jannoColumn COLNAME</h1>	List additional fields from the janno files, using 700 the .janno column heading name, such as "Country", 701 "Site", "Date_C14_Uncal_BP", etc..

702 **--raw** Return the output table as tab-separated values
703 without header. This is useful for piping into grep
704 or awk.

705 **--onlyLatest** Consider only the latest versions of packages, or the
706 groups and individuals within the latest versions of
707 packages, respectively.

708 To list packages from your local repositories, as seen above you can run

709 `trident list -d ... -d ... --packages`

710 This will yield a nicely formatted table of all packages, their version and the number of individuals in them.

711 You can use **--remote** to show packages on the remote server. For example

712 `trident list --packages --remote --archive "community-archive"`

713 will result in a view of all packages available in one of the Poseidon public archives. Just as for **fetch**, the
714 **--archive** flag allows to choose which public archive to query.

715 Independent of whether you query a local or an online archive, you can not just list packages, but also groups,
716 as defined in the third column of EIGENSTRAT `.ind` files (or the first/last column of a PLINK `.fam` file), and
717 individuals with the flags **--groups** and **--individuals** (instead of **--packages**).

718 The **--individuals** flag additionally provides a way to immediately access information from `.janno` files
719 on the command line. This works with the `-j/--jannoColumn` option. For example adding `-j Country -j`
720 `Date_C14_Uncal_BP` to the commands above will add the `Country` and the `Date_C14_Uncal_BP` columns to the
721 respective output tables.

722 Note that if you want a less fancy table, for example because you want to load this into Excel, or pipe into
723 another command that cannot deal with the table layout, you can use the **--raw** option to output that table as
724 a simple tab-delimited stream.

725 4.2 Summarise command

726 **summarise** prints some general summary statistics for a given poseidon dataset taken from the `.janno` files.

727 Command line details

728 Usage: `trident summarise (-d|--baseDir DIR) [--raw]`

729

730 Get an overview over the content of one or multiple Poseidon packages

731

732 Available options:

733 **-h,--help** Show this help text

734 **-d,--baseDir DIR** A base directory to search for Poseidon packages.

735 **--raw** Return the output table as tab-separated values
736 without header. This is useful for piping into grep
737 or awk.

738 You can run it with

739 `trident summarise -d ... -d ...`

740 which will show you context information like – among others – the number of individuals in the dataset, their
 741 sex distribution, the mean age of the samples (for ancient data) or the mean coverage on the 1240K SNP array
 742 in a table. `summarise` depends on complete `.janno` files and will silently ignore missing information.
 743 You can use the `--raw` option to output the summary table in a simple, tab-delimited layout.

744 4.3 Survey command

745 `survey` tries to indicate package completeness (mostly focused on `.janno` files) for poseidon datasets.

746 Command line details

747 Usage: `trident survey (-d|--baseDir DIR) [--raw] [--onlyLatest]`

749 Survey the degree of context information completeness for Poseidon packages

751 Available options:

752	<code>-h,--help</code>	Show this help text
753	<code>-d,--baseDir DIR</code>	A base directory to search for Poseidon packages.
754	<code>--raw</code>	Return the output table as tab-separated values without header. This is useful for piping into <code>grep</code> or <code>awk</code> .
757	<code>--onlyLatest</code>	Consider only the latest versions of packages, or the groups and individuals within the latest versions of packages, respectively.

760 Running

761 `trident survey -d ... -d ...`

762 will yield a table with one row for each package. See `trident survey -h` for a legend which cell of this table
 763 means what.

764 Again you can use the `--raw` option to output the survey table in a tab-delimited format.

765 4.4 Validate command

766 `validate` checks Poseidon packages and individual package components for structural correctness.

767 Command line details

768 Usage: `trident validate ((-d|--baseDir DIR) [--ignoreGeno] [--fullGeno]
 769 [--ignoreDuplicates] [-c|--ignoreChecksums]
 770 [--ignorePoseidonVersion] |
 771 --pyml FILE | (-p|--genoOne FILE) | --inFormat FORMAT
 772 --genoFile FILE --snpFile FILE --indFile FILE |
 773 --janno FILE | --ssf FILE | --bib FILE) [--noExitCode]
 774 [--onlyLatest]`

776 Check Poseidon packages or package components for structural correctness

778 Available options:

```

779  -h,--help                Show this help text
780  -d,--baseDir DIR        A base directory to search for Poseidon packages.
781  --ignoreGeno            Ignore snp and geno file.
782  --fullGeno              Test parsing of all SNPs (by default only the first
783                          100 SNPs are probed).
784  --ignoreDuplicates      Do not stop on duplicated individual names in the
785                          package collection.
786  -c,--ignoreChecksums    Whether to ignore checksums. Useful for speedup in
787                          debugging.
788  --ignorePoseidonVersion Read packages even if their poseidonVersion is not
789                          compatible with trident.
790  --pyml FILE             Path to a POSEIDON.yml file.
791  -p,--genoOne FILE       One of the input genotype data files. Expects .bed,
792                          .bim or .fam for PLINK and .geno, .snp or .ind for
793                          EIGENSTRAT. The other files must be in the same
794                          directory and must have the same base name.
795  --inFormat FORMAT       The format of the input genotype data: EIGENSTRAT or
796                          PLINK. Only necessary for data input with --genoFile
797                          + --snpFile + --indFile.
798  --genoFile FILE         Path to the input geno file.
799  --snpFile FILE          Path to the input snp file.
800  --indFile FILE          Path to the input ind file.
801  --janno FILE            Path to a .janno file.
802  --ssf FILE              Path to a .ssf file.
803  --bib FILE              Path to a .bib file.
804  --noExitCode            Do not produce an explicit exit code.
805  --onlyLatest            Consider only the latest versions of packages, or the
806                          groups and individuals within the latest versions of
807                          packages, respectively.
808  You can run it with
809  trident validate -d ... -d ...
810  to check packages and it will either report a success (Validation passed) or failure with specific error messages.
811  Instead of validating entire packages with -d you can also apply it to individual files and package com-
812  ponents: --pyml (POSEIDON.yml), -p | --inFormat + --genoFile + --snpFile + --indFile (genotype
813  data), --janno (.janno file), --ssf (.ssf file) or --bib (.bib file). In this case validate attempts to read and
814  parse the respective files individually and reports any issues it encounters. Note that this considers the files in
815  isolation and does not include any cross-file consistency checks.
816  When applied to packages, validate tries to ensure that each package adheres to the Poseidon package
817  specification. Here is a list of what is checked:
818  • Structural correctness of the POSEIDON.yml file.
819  • Presence of all files references in the POSEIDON.yml file.
820  • Full structural correctness of .janno, .ssf and .bib file.
821  • Superficial correctness of genotype data files by parsing the first 100 SNPs. A full check that parses all

```

822 SNPs can be triggered with the `--fullGeno` option. `--ignoreGeno`, on the other hand, causes `validate`
823 to ignore the genotype data entirely, which speeds up the validation significantly.

- 824 • Correspondence of BibTeX keys in `.bib` and `.janno`
- 825 • Correspondence of sample IDs in `.janno` and `.ssf`.
- 826 • Correspondence of sample and group IDs in `.janno` and genotype data files.

827 In fact much of this validation already runs as part of the general package reading pipeline invoked for other
828 trident subcommands (e.g. `forge`). `validate` is meant to be more thorough/brittle, though, and will explicitly
829 fail if even a single package is broken. For special cases more flexibility can be enabled with the options
830 `--ignoreDuplicates`, `--ignoreChecksums` and `--ignorePoseidonVersion`.

831 Remember to run `validate` it with `--debug` to get more information in case the default output is not sufficient
832 to analyse an issue.