

Guide for trident v1.1.10.2

Contents

1	The trident CLI	1
1.1	General notes	4
1.1.1	Logging and command line output	4
1.1.2	Duplicates	4
1.1.3	Group names in .fam files	4
1.1.4	Whitespaces in the .janno file	4
2	Package creation and manipulation commands	4
2.1	Init command	4
2.2	Fetch command	6
2.3	Forge command	7
2.3.1	The forge selection language	10
2.3.2	Treatment of the .janno file while merging	11
2.3.3	Other options	12
2.4	Genoconvert command	12
2.5	Update command	14
3	Inspection commands	16
3.1	List command	16
3.2	Summarise command	17
3.3	Survey command	17
3.4	Validate command	18

1 The trident CLI

Trident is a command line software tool structured in multiple subcommands. If you installed it properly you can call it on the command line by typing `trident`. This will show an overview of the general options and all subcommands, which are explained in detail below.

Usage: `trident [--version] [--logMode ARG] [--errLength ARG]`

`[--inPlinkPopName ARG] (COMMAND | COMMAND)`

`trident` is a management and analysis tool for Poseidon packages. Report issues here: <https://github.com/poseidon-framework/poseidon-hs/issues>

Available options:

```

34  -h,--help          Show this help text
35  --version          Show version number
36  --logMode ARG      How information should be reported: NoLog, SimpleLog,
37                    DefaultLog, ServerLog or VerboseLog
38                    (default: DefaultLog)
39  --errLength ARG    After how many characters should a potential error
40                    message be truncated. "Inf" for no truncation.
41                    (default: CharCount 1500)
42  --inPlinkPopName ARG Where to read the population/group name from the FAM
43                    file in Plink-format. Three options are possible:
44                    asFamily (default) | asPhenotype | asBoth.
45
46  Package creation and manipulation commands:
47  init              Create a new Poseidon package from genotype data
48  fetch             Download data from a remote Poseidon repository
49  forge             Select packages, groups or individuals and create a
50                    new Poseidon package from them
51  genoconvert       Convert the genotype data in a Poseidon package to a
52                    different file format
53  update            Update POSEIDON.yml files automatically
54
55  Inspection commands:
56  list              List packages, groups or individuals from local or
57                    remote Poseidon repositories
58  summarise         Get an overview over the content of one or multiple
59                    Poseidon packages
60  summarize         Synonym for summarise
61  survey            Survey the degree of context information completeness
62                    for Poseidon packages
63  validate          Check one or multiple Poseidon packages for
64                    structural correctness
65
66  Trident allows to work directly with genotype data (see -p below), but its optimized for the interaction
67  with Poseidon packages, which wrap and contextualize the data. Most trident subcommands therefore have a
68  central parameter, called --baseDir or simply -d to specify one or more base directories to look for packages.
69  For example, if all Poseidon packages live inside a repository at /path/to/poseidon/packages you would
70  simply say trident <subcommand> -d /path/to/poseidon/dirs/ and trident would automatically search
71  all subdirectories inside of the repository for valid Poseidon packages (as identified by valid POSEIDON.yml
72  files).
73
74  You can arrange a poseidon repository in a hierarchical way. For example:
75
76  /path/to/poseidon/packages
77      /modern
78          /2019_poseidon_package1
79          /2019_poseidon_package2
80      /ancient

```

```

78     /...
79     /...
80 /Reference_Genomes
81     /...
82     /...

```

83 You can use this structure to select only the level of packages you're interested in, even individual ones, and you
84 can make use of the fact that `-d` can be given multiple times.

85 Being able to specify one or multiple repositories is often not enough, as you may have your own data to
86 co-analyse with the main repository. This is easy to do, as you simply need to provide your own genotype data
87 as yet another Poseidon package to be added to your `trident` command. For example, let's say you have
88 genotype data in `EIGENSTRAT` format (`trident` supports `EIGENSTRAT` and `PLINK` as formats.):

```

89 ~/my_project/my_project.geno
90 ~/my_project/my_project.snp
91 ~/my_project/my_project.ind

```

92 then you can make that to a skeleton Poseidon package with the `init` command. You can also do it manually
93 by simply adding a `POSEIDON.yml` file, with for example the following content:

```

94 poseidonVersion: 2.5.0
95 title: My_awesome_project
96 description: Unpublished genetic data from my awesome project
97 contributor:
98   - name: Stephan Schiffels
99     email: schiffels@institute.org
100 packageVersion: 0.1.0
101 lastModified: 2020-10-07
102 genotypeData:
103   format: EIGENSTRAT
104   genoFile: my_project.geno
105   snpFile: my_project.snp
106   indFile: my_project.ind
107 jannoFile: my_project.janno
108 bibFile: sources.bib

```

109 Two remarks: 1) all file paths are considered *relative* to the directory in which `POSEIDON.yml` resides. Here we
110 assume that you put this file into the same directory as the three genotype files. 2) Besides the genotype data
111 files there are two (technically optional) files referenced by this example `POSEIDON.yml` file: `sources.bib`
112 and `my_project.janno`. Of course you can add them manually - `init` automatically creates empty dummy
113 versions.

114 Once you have set up your own “Poseidon” package (which is really only a skeleton so far), you can add it to
115 your `trident` analysis, by simply adding your project directory to the command using `-d`, for example:

```

116 trident list -d /path/to/poseidon/packages/modern \
117   -d /path/to/poseidon/packages/ReferenceGenomes
118   -d ~/my_project --packages

```

1.1 General notes

1.1.1 Logging and command line output

For all subcommands the general argument `--logMode` defines how trident reports messages (to stderr) on the command line:

- *NoLog*: Hides all messages.
- *SimpleLog*: Plain and simple output to stderr.
- *DefaultLog*: Adds severity indicators before each message. (default setting)
- *ServerLog*: Additionally adds timestamps before each message.
- *VerboseLog*: Shows not just messages on the log levels `Info`, `Warning` and `Error` like the other modes, but also on the more verbose level `Debug`. Use this for debugging.

1.1.2 Duplicates

- If multiple packages in a package repository share the same `title`, then trident will try to select the one with the highest version number. If this is not sufficient to resolve the conflict, trident will stop.
- Individual/sample names (`Poseidon_ID`s) within one package have to be unique, or trident will stop.
- We generally also discourage ID duplicates across packages in package repositories, but trident will generally continue with them after printing a warning. This does not apply for `validate`, by default (you can change this behaviour with `--ignoreDuplicates`), and `forge`. `forge` offers a special mechanism to resolve duplicates within its selection language (see below).

1.1.3 Group names in .fam files

The `.fam` file of Plink-formatted genotype data is used inconsistently across different popular aDNA software tools to store group/population name information. The (global) option `--inPlinkPopName` with the arguments `asFamily` (default), `asPhenotype` and `asBoth` allows to control the reading of the population name from Plink `.fam` files. The subcommands that write genotype data (`forge`, `genoconvert`) have a corresponding option `--outPlinkPopName` to specify this for the output.

1.1.4 Whitespaces in the .janno file

While reading the `.janno` file `trident` trims all leading and trailing whitespaces around individual cells. Also all instances of the `No-Break Space` unicode character will be removed. This means these whitespaces will not be preserved when a package is `forge`d.

2 Package creation and manipulation commands

2.1 Init command

`init` creates a new, valid Poseidon package from genotype data files. It adds a valid `POSEIDON.yml` file, a dummy `.janno` file for context information and an empty `.bib` file for literature references.

[Click here for command line details](#)

```
Usage: trident init ((-p|--genoOne ARG) | --inFormat ARG --genoFile ARG
                    --snpFile ARG --indFile ARG) [--snpSet ARG]
                    (-o|--outPackagePath ARG) [-n|--outPackageName ARG]
```

```

155             [--minimal]
156     Create a new Poseidon package from genotype data
157
158     Available options:
159     -h,--help                Show this help text
160     -p,--genoOne ARG         one of the input genotype data files. Expects .bed or
161                               .bim or .fam for PLINK and .geno or .snp or .ind for
162                               EIGENSTRAT. The other files must be in the same
163                               directory and must have the same base name
164     --inFormat ARG           the format of the input genotype data: EIGENSTRAT or
165                               PLINK (only necessary for data input with --genoFile
166                               + --snpFile + --indFile)
167     --genoFile ARG           the input geno file path
168     --snpFile ARG            the input snp file path
169     --indFile ARG            the input ind file path
170     --snpSet ARG             the snpSet of the package: 1240K, HumanOrigins or
171                               Other. (only relevant for data input with
172                               -p|--genoOne or --genoFile + --snpFile + --indFile,
173                               because the packages in a -d|--baseDir already have
174                               this information in their respective POSEIDON.yml
175                               files) Default: Other
176     -o,--outPackagePath ARG  the output package directory path
177     -n,--outPackageName ARG  the output package name - this is optional: If no
178                               name is provided, then the package name defaults to
179                               the basename of the (mandatory) --outPackagePath
180                               argument
181     --minimal                should only a minimal output package be created?
182
183     The command
184
185     trident init \
186     --inFormat EIGENSTRAT/PLINK \
187     --genoFile path/to/geno_file \
188     --snpFile path/to/snp_file \
189     --indFile path/to/ind_file \
190     --snpSet 1240K|HumanOrigins|Other \
191     -o path/to/new_package_name
192
193     requires the format ( --inFormat ) of your input data (either EIGENSTRAT or PLINK), the paths to the
194     respective files ( --genoFile , --snpFile , --indFile ), and optionally the “shape” of these files ( --snpSet ),
195     so if they cover the 1240K, the HumanOrigins or an Other SNP set. A simpler interface added in trident
196     0.29.0 is available with -p (+ --snpSet) .

```

	EIGENSTRAT	PLINK
genoFile	.geno	.bed
snpFile	.snp	.bim
indFile	.ind	.fam

194 The output package of `init` is created as a new directory `-o`, which should not already exist, and gets the
195 package `title` corresponding to the basename of `-o`. You can also set the title explicitly with `-n`. The
196 `--minimal` flag causes `init` to create a minimal package with a very basic `POSEIDON.yml` and no `.bib` and
197 `.janno` files.

198 2.2 Fetch command

199 `fetch` allows to download Poseidon packages from a remote Poseidon server. Read more about this repository
200 [here](#).

201 [Click here](#) for command line details

```
202 Usage: trident fetch (-d|--baseDir DIR)
203             (--downloadAll |
204             (--fetchFile ARG | (-f|--fetchString ARG)))
205             [--remoteURL ARG] [-u|--upgrade]
```

206 Download data from a remote Poseidon repository

207
208 Available options:

209	<code>-h,--help</code>	Show this help text
210	<code>-d,--baseDir DIR</code>	a base directory to search for Poseidon Packages (could be a Poseidon repository)
211	<code>--downloadAll</code>	download all packages the server is offering
212	<code>--fetchFile ARG</code>	A file with a list of packages. Works just as <code>-f</code> , but multiple values can also be separated by newline, not just by comma. <code>-f</code> and <code>--fetchFile</code> can be combined.
213	<code>-f,--fetchString ARG</code>	List of packages to be downloaded from the remote server. Package names should be wrapped in asterisks: <code>*package_title*</code> . You can combine multiple values with comma, so for example: <code>"*package_1*, *package_2*, 214 *package_3*"</code> . <code>fetchString</code> uses the same parser as 215 <code>forgeString</code> , but does not allow excludes. If groups or individuals are specified, then packages which include these groups or individuals are included in the download.
216	<code>--remoteURL ARG</code>	URL of the remote Poseidon server (default: "https://c107-224.cloud.gwdg.de")
217	<code>-u,--upgrade</code>	overwrite outdated local package versions

218 It works with

```
219 trident fetch -d ... -d ... \  
220 -f "*package_title_1*,*package_title_2*,*package_title_3*,group_name,<Individual1>"
```

221 and the entities you want to download must be listed either in a simple string of comma-separated values, which
222 can be passed via `-f` / `--fetchString`, or in a text file (`--fetchFile`). Entities are then combined from
223 these sources.

224 Entities are specified using a special syntax (see also the documentation of `forge` below): Package titles

are wrapped in asterisks: *package_title*, group names are spelled as is, and individual names are wrapped in angular brackets, like `<Individual1>`. Fetch will figure out which packages need to be downloaded to include all specified entities. `--downloadAll`, which can be given instead of `-f` and `--fetchFile`, causes fetch to download all packages from the server. The downloaded packages are added in the first (!) `-d` directory (which gets created if it doesn't exist), but downloads are only performed if the respective packages are not already present in an up-to-date version in any of the `-d` dirs.

Note that `trident fetch` makes most sense in combination with `trident list --remote`: First one can inspect what is available on the server, then one can create a custom fetch command.

`fetch` also has the optional arguments `--remote https://..."` to name an alternative poseidon server. The default points to the **DAG server**.

To overwrite outdated package versions with `fetch`, the `-u` / `--upgrade` flag has to be set. Note that many file systems do not offer a way to recover overwritten files. So be careful with this switch.

2.3 Forge command

`forge` creates new Poseidon packages by extracting and merging packages, populations and individuals from your Poseidon repositories.

Click here for command line details

```
Usage: trident forge ((-d|--baseDir DIR) |
                    ((-p|--genoOne ARG) | --inFormat ARG --genoFile ARG
                    --snpFile ARG --indFile ARG) [--snpSet ARG])
                    [--forgeFile ARG | (-f|--forgeString ARG)]
                    [--selectSnps ARG] [--intersect] [--outFormat ARG]
                    [--minimal] [--onlyGeno] (-o|--outPackagePath ARG)
                    [-n|--outPackageName ARG] [--packagewise]
                    [--outPlinkPopName ARG]
```

Select packages, groups or individuals and create a new Poseidon package from them

Available options:

<code>-h,--help</code>	Show this help text
<code>-d,--baseDir DIR</code>	a base directory to search for Poseidon Packages (could be a Poseidon repository)
<code>-p,--genoOne ARG</code>	one of the input genotype data files. Expects .bed or .bim or .fam for PLINK and .geno or .snp or .ind for EIGENSTRAT. The other files must be in the same directory and must have the same base name
<code>--inFormat ARG</code>	the format of the input genotype data: EIGENSTRAT or PLINK (only necessary for data input with <code>--genoFile</code> + <code>--snpFile</code> + <code>--indFile</code>)
<code>--genoFile ARG</code>	the input geno file path
<code>--snpFile ARG</code>	the input snp file path
<code>--indFile ARG</code>	the input ind file path
<code>--snpSet ARG</code>	the snpSet of the package: 1240K, HumanOrigins or

277 Other. (only relevant for data input with
278 -p|--genoOne or --genoFile + --snpFile + --indFile,
279 because the packages in a -d|--baseDir already have
280 this information in their respective POSEIDON.yml
281 files) Default: Other

282 --forgeFile ARG A file with a list of packages, groups or individual
283 samples. Works just as -f, but multiple values can
284 also be separated by newline, not just by comma.
285 Empty lines are ignored and comments start with "#",
286 so everything after "#" is ignored in one line.
287 Multiple instances of -f and --forgeFile can be
288 given. They will be evaluated according to their
289 input order on the command line.

290 -f,--forgeString ARG List of packages, groups or individual samples to be
291 combined in the output package. Packages follow the
292 syntax *package_title*, populations/groups are simply
293 group_id and individuals <individual_id>. You can
294 combine multiple values with comma, so for example:
295 "*package_1*, <individual_1>, <individual_2>,"
296 group_1". Duplicates are treated as one entry.
297 Negative selection is possible by prepending "-" to
298 the entity you want to exclude (e.g. "*package_1*,
299 -<individual_1>, -group_1"). forge will apply
300 excludes and includes in order. If the first entity
301 is negative, then forge will assume you want to merge
302 all individuals in the packages found in the baseDirs
303 (except the ones explicitly excluded) before the
304 exclude entities are applied. An empty forgeString
305 (and no --forgeFile) will therefore merge all
306 available individuals. If there are individuals in
307 your input packages with equal individual id, but
308 different main group or source package, they can be
309 specified with the special syntax
310 "<package:group:individual>".

311 --selectSnps ARG To extract specific SNPs during this forge operation,
312 provide a Snp file. Can be either Eigenstrat (file
313 ending must be '.snp') or Plink (file ending must be
314 '.bim'). When this option is set, the output package
315 will have exactly the SNPs listed in this file. Any
316 SNP not listed in the file will be excluded. If
317 option '--intersect' is also set, only the SNPs
318 overlapping between the SNP file and the forged
319 packages are output.

320 --intersect Whether to output the intersection of the genotype
321 files to be forged. The default (if this option is

not set) is to output the union of all SNPs, with genotypes defined as missing in those packages which do not have a SNP that is present in another package. With this option set, the forged dataset will typically have fewer SNPs, but less missingness.

`--outFormat ARG` the format of the output genotype data: EIGENSTRAT or PLINK. Default: PLINK

`--minimal` should only a minimal output package be created?

`--onlyGeno` should only the resulting genotype data be returned? This means the output will not be a Poseidon package

`-o,--outPackagePath ARG` the output package directory path

`-n,--outPackageName ARG` the output package name - this is optional: If no name is provided, then the package name defaults to the basename of the (mandatory) `--outPackagePath` argument

`--packagewise` Skip the within-package selection step in forge. This will result in outputting all individuals in the relevant packages, and hence a superset of the requested individuals/groups. It may result in better performance in cases where one wants to forge entire packages or almost entire packages. Details: Forge conceptually performs two types of selection: First, it identifies which packages in the supplied base directories are relevant to the requested forge, i.e. whether they are either explicitly listed using `*PackageName*`, or because they contain selected individuals or groups. Second, within each relevant package, individuals which are not requested are removed. This option skips only the second step, but still performs the first.

`--outPlinkPopName ARG` Where to write the population/group name into the FAM file in Plink-format. Three options are possible: `asFamily` (default) | `asPhenotype` | `asBoth`. See also `--inPlinkPopName`.

`forge` can be used with

```
trident forge -d ... -d ... \
  -f "*package_name*, group_id, <individual_id>" \
  -o path/to/new_package_name
```

where the entities (packages, groups/populations, individuals/samples) you want in the output package can be denoted either as a string on the command line (`-f / --forgeString`), or in an input text file (`--forgeFile`). See the section below for the syntax of this selection language. Do not forget to wrap the `--forgeString` query in quotes.

Including one or multiple Poseidon packages with `-d` is not the only way to include data for a forge operation. It is also possible to consider unpackaged genotype data directly with `-p (+ --snpSet)` or

366 `--inFormat + --genoFile + --snpFile + --indFile (+ --snpSet)`. This makes the following example
367 possible, where we merge data from one Poseidon package and two genotype datasets to get a new EIGENSTRAT
368 dataset.

```
369 trident forge \  
370   -d 2017_GonzalesFortesCurrentBiology \  
371   -p 2018_VeeramahPNAS/2018_VeeramahPNAS.fam \  
372   --inFormat PLINK \  
373   --genoFile 2017_HaberAJHG/2017_HaberAJHG.bed \  
374   --snpFile 2017_HaberAJHG/2017_HaberAJHG.bim \  
375   --indFile 2017_HaberAJHG/2017_HaberAJHG.fam \  
376   -f "<STR241.SG>,<ERS1790729.SG>,Iberia_HG.SG" \  
377   -o testpackage \  
378   --outFormat EIGENSTRAT \  
379   --onlyGeno
```

380 2.3.1 The forge selection language

381 The text in `--forgeString` and `--forgeFile` are parsed as a domain specific query language that describes
382 precisely which entities should be compiled in the output package of a given `forge` operation. The language
383 has multiple syntactic elements and a specific evaluation logic.

384 In general a `--forgeString` query consists of multiple entities, separated by `,`. The main entities are Poseidon
385 packages, groups/populations and individuals/samples:

- 386 • Each package title is surrounded by `* : *package*`. That means if you want all individuals of the Poseidon
387 package `2019_Jeong_InnerEurasia` in the output package you would add `*2019_Jeong_InnerEurasia*`
388 to the query.
- 389 • Groups/populations are not specially marked: `group`. So to get all individuals of the group
390 `Swiss_Roman_period`, you would simply add `Swiss_Roman_period`.
- 391 • Individuals/samples are surrounded by `<` and `>`: `<individual>`. `ALA026` therefore becomes
392 `<ALA026>`. A second way to denote individuals is with the more verbose and specific syntax
393 `<package:group:individual>`. Such defined individuals take precedence over differently defined ones
394 (so: directly with `<individual>` or as a subset of `*package*` or `group`). This allows to resolve
395 duplication issues precisely – at least in cases where the duplicated individuals differ in source package or
396 primary group.

397 In the `--forgeFile` each line is treated as a separate `forgeString`, empty lines are ignored and `#`s start
398 comments. So this is a valid `forgeFile`:

```
399 # Packages  
400 *package1*, *package2*  
401  
402 # Groups and individuals from other packages beyond package1 and package2  
403 group1, <individual1>, group2, <individual2>, <individual3>  
404  
405 # group2 has two outlier individuals that should be ignored  
406 -<bad_individual1> # This one has very low coverage  
407 -<bad_individual2> # This one is from a different time period
```

By prepending `-` to the bad individuals, we can exclude them from the forged package. `forge` figures out the final list of samples to include by executing all forge-entities in order. So an entity list `*PackageA*,-<Individual1>,GroupA` may result in a different outcome than `*PackageA*,GroupA,-<Individual1>`, depending on whether `<Individual1>` belongs to `GroupA` or not. If the forge entity list starts with a negative entity, or if the entity list is empty, `forge` will implicitly assume you want to include all individuals in all packages found in the baseDirs (except the ones explicitly excluded, of course).

An empty forgeString will therefore merge all available individuals.

2.3.2 Treatment of the .janno file while merging

`forge` merges and subsets .janno files along with the genotype data. If a package lacks a .janno file, then a basic one will be created internally based on the information in the genotype data, and used for the output. Missing columns across packages will be filled with `n/a`.

For merging two .janno files **A** and **B** the following rules apply regarding undefined, arbitrary additional columns:

- If **A** has an additional column which is not in **B** then empty cells in the rows imported from **B** are filled with `n/a`.
- If **A** and **B** share additional columns with identical column name, then they are treated as semantically identical units and merged accordingly.
- In the resulting .janno file, all additional columns from both **A** and **B** are sorted alphabetically and appended after the normal, specified variables.

The following example illustrates the described behaviour:

A.janno

Poseidon_ID	Group_Name	Genetic_Sex	AdditionalColumn1	AdditionalColumn2
XXX011	POP1	M	A	D
XXX012	POP2	F	B	E
XXX013	POP1	M	C	F

B.janno

Poseidon_ID	Group_Name	Genetic_Sex	AdditionalColumn3	AdditionalColumn2
YYY022	POP5	F	G	J
YYY023	POP5	F	H	K
YYY024	POP5	M	I	L

A.janno + B.janno

Poseidon_ID	Group_Name	Genetic_Sex	AdditionalColumn1	AdditionalColumn2	AdditionalColumn3
XXX011	POP1	M	A	D	n/a
XXX012	POP2	F	B	E	n/a
XXX013	POP1	M	C	F	n/a
YYY022	POP5	F	n/a	J	G

Poseidon_ID	Group_Name	Genetic_Sex	AdditionalColumn1	AdditionalColumn2	AdditionalColumn3
YYY023	POP5	F	n/a	K	H
YYY024	POP5	M	n/a	L	I

2.3.3 Other options

Just as for `init` the output package of `forge` is created as a new directory `-o`. The title can also be explicitly defined with `-n`.

`--minimal` allows for the creation of a minimal output package without `.bib` and `.janno`. This is especially useful for data analysis pipelines, where only the genotype data is required. Even more basic output comes with `--onlyGeno`, which means that only the genotype data is returned without any Poseidon package.

`forge` has a an optional flag `--intersect`, that defines, if the genotype data from different packages should be merged with an **union** or an **intersect** operation. The default (if this option is not set) is to output the union of all SNPs, with genotypes defined as missing in samples from packages which do not have a SNP that is present in another package. With this option set, on the other hand, the forged dataset will typically have fewer SNPs, but less missingness.

`--intersect` also influences the automatic determination of the `snpSet` field in the POSEIDON.yml file for the resulting package. If the `snpSet` s of all input packages are identical, then the resulting package will just inherit this configuration. Otherwise `forge` applies the following pairwise merging logic:

Input snpSet A	Input snpSet B	<code>--intersect</code>	Ouput snpSet
Other	*	*	Other
1240K	HumanOrigins	True	HumanOrigins
1240K	HumanOrigins	False	1240K

`--selectSnps` allows to provide `forge` with a SNP file in EIGENSTRAT (`.snp`) or PLINK (`.bim`) format to create a package with a specific selection. When this option is set, the output package will have exactly the SNPs listed in this file. Any SNP not listed in the file will be excluded. If `--intersect` is also set, only the SNPs overlapping between the SNP file and the forged packages are output.

Merging genotype data across different data sources and file formats is tricky. `forge` is more verbose about potential issues, if the `--logMode` flag is set to `VerboseLog`.

The `--onlyGeno` command specifies that only genotype data should be output, not an entire Poseidon package.

With `--packagewise` the within-package selection step in `forge` can be skipped. This will result in outputting all individuals in the relevant packages, and hence a superset of the requested individuals/groups. It may result in better performance in cases where one wants to forge entire packages.

2.4 Genoconvert command

`genoconvert` converts the genotype data in a Poseidon package to a different file format. The respective entries in the POSEIDON.yml file are changed accordingly.

[Click here for command line details](#)

```

458 Usage: trident genoconvert ((-d|--baseDir DIR) |
459     ((-p|--genoOne ARG) | --inFormat ARG --genoFile ARG
460     --snpFile ARG --indFile ARG) [--snpSet ARG])
461     --outFormat ARG [--onlyGeno]
462     [-o|--outPackagePath ARG] [--removeOld]
463     [--outPlinkPopName ARG]
464     Convert the genotype data in a Poseidon package to a different file format
465
466 Available options:
467     -h,--help                Show this help text
468     -d,--baseDir DIR         a base directory to search for Poseidon Packages
469                               (could be a Poseidon repository)
470     -p,--genoOne ARG         one of the input genotype data files. Expects .bed or
471                               .bim or .fam for PLINK and .geno or .snp or .ind for
472                               EIGENSTRAT. The other files must be in the same
473                               directory and must have the same base name
474     --inFormat ARG           the format of the input genotype data: EIGENSTRAT or
475                               PLINK (only necessary for data input with --genoFile
476                               + --snpFile + --indFile)
477     --genoFile ARG           the input geno file path
478     --snpFile ARG            the input snp file path
479     --indFile ARG           the input ind file path
480     --snpSet ARG            the snpSet of the package: 1240K, HumanOrigins or
481                               Other. (only relevant for data input with
482                               -p|--genoOne or --genoFile + --snpFile + --indFile,
483                               because the packages in a -d|--baseDir already have
484                               this information in their respective POSEIDON.yml
485                               files) Default: Other
486     --outFormat ARG         the format of the output genotype data: EIGENSTRAT or
487                               PLINK.
488     --onlyGeno              should only the resulting genotype data be returned?
489                               This means the output will not be a Poseidon package
490     -o,--outPackagePath ARG the output package directory path - this is optional:
491                               If no path is provided, then the output is written to
492                               the directories where the input genotype data file
493                               (.bed/.geno) is stored
494     --removeOld             Remove the old genotype files when creating the new
495                               ones
496     --outPlinkPopName ARG   Where to write the population/group name into the FAM
497                               file in Plink-format. Three options are possible:
498                               asFamily (default) | asPhenotype | asBoth. See also
499                               --inPlinkPopName.
500
501 With the default setting
502 trident genoconvert -d ... -d ... --outFormat EIGENSTRAT|PLINK

```

all packages in `-d` will be converted to the desired `--outFormat` (either `EIGENSTRAT` or `PLINK`), if the data is not already in this format. This includes updating the respective `POSEIDON.yml` files.

The “old” data is not deleted, but kept around. That means conversion can result in a package with both `PLINK` and `EIGENSTRAT` data, but only one is linked in the `POSEIDON.yml` file, and that is what will be used by `trident`. To delete the old data in the conversion you can add the `--removeOld` flag.

Instead of `-d` to change Poseidon packages, the `-p (+ --snpSet)` or `--inFormat + --genoFile + --snpFile + --indFile` allow to directly convert genotype data that is not wrapped in a Poseidon package and store it to a directory given in `-o`. See this example:

```
trident genoconvert \  
-p 2018_Mittnik_Baltic/Mittnik_Baltic.bed \  
--outFormat EIGENSTRAT \  
-o my_directory
```

2.5 Update command

`update` automatically harmonizes `POSEIDON.yml` files of one or multiple packages if the packages were changed. This is not an automatic update from one Poseidon version to the next!

Click here for command line details

```
Usage: trident update (-d|--baseDir DIR) [--poseidonVersion ARG]   
                [--ignorePoseidonVersion] [--versionComponent ARG]   
                [--noChecksumUpdate] [--newContributors ARG]   
                [--logText ARG] [--force]
```

Update `POSEIDON.yml` files automatically

Available options:

<code>-h,--help</code>	Show this help text
<code>-d,--baseDir DIR</code>	a base directory to search for Poseidon Packages (could be a Poseidon repository)
<code>--poseidonVersion ARG</code>	Poseidon version the packages should be updated to: e.g. "2.5.3" (default: Nothing)
<code>--ignorePoseidonVersion</code>	Read packages even if their poseidonVersion is not compatible with the trident version. The assumption is, that the package is already structurally adjusted to the trident version and only the version number is lagging behind.
<code>--versionComponent ARG</code>	Part of the package version number in the <code>POSEIDON.yml</code> file that should be updated: Major, Minor or Patch (see https://semver.org) (default: Patch)
<code>--noChecksumUpdate</code>	Should update of checksums in the <code>POSEIDON.yml</code> file be skipped
<code>--ignoreGeno</code>	ignore SNP and <code>GenoFile</code>
<code>--newContributors ARG</code>	Contributors to add to the <code>POSEIDON.yml</code> file in the form "[Firstname Lastname](Email address);..."

```

544 --logText ARG          Log text for this version jump in the CHANGELOG file
545                        (default: "not specified")
546 --force                Normally the POSEIDON.yml files are only changed if
547                        the poseidonVersion is adjusted or any of the
548                        checksums change. With --force a package version
549                        update can be triggered even if this is not the case.

```

550 It can be called with a lot of optional arguments

```

551 trident update -d ... -d ... \
552   --poseidonVersion "X.X.X" \
553   --versionComponent Major/Minor/Patch \
554   --noChecksumUpdate
555   --ignoreGeno
556   --newContributors "[Firstname Lastname](Email address);..."
557   --logText "short description of the update"
558   --force

```

559 By default `update` will not edit a package's POSEIDON.yml file, even when arguments like `--versionComponent`,
560 `--newContributors` or `--logText` are explicitly set. This default exists to run the function on a large set of
561 packages where only few of them were edited and need an active update. A package will only be modified by
562 `update` if either

- 563 • any of the files with checksums (e.g. the genotype data) in it were modified,
- 564 • the `--poseidonVersion` argument differs from the `poseidonVersion` in the package's POSEIDON.yml
565 file
- 566 • or the `--force` flag was set in `update`.

567 If any of these applies to a package in the search directory (`--baseDir / -d`), it will be updated. This includes
568 the following steps:

- 569 • If `--poseidonVersion` is different from the `poseidonVersion` field in the package, then that will be
570 updated.
- 571 • The `packageVersion` will be incremented. If `--versionComponent` is not set, then it falls back to
572 `Patch`, so a change in the last position of the three digit version number. `Minor` increments the middle,
573 and `Major` the first position (see [semantic versioning](#)).
- 574 • The `lastModified` field will be updated to the current day (based on your computer's system time).
- 575 • The contributors in `--newContributors` will be added to the `contributor` field if they're not there
576 already.
- 577 • If any checksums changed, then they will be updated. If certain checksums are not set yet, then they will
578 be added. The checksum update can be skipped with `--noChecksumUpdate` or partially skipped for the
579 genotype data with `--ignoreGeno`.
- 580 • The CHANGELOG.md file will be updated with a new row for the new version and the text in `--logText`
581 (default: "not specified"), which will be appended as the first line of the file. If no CHANGELOG.md file
582 exists, then it will be created and referenced in the POSEIDON.yml file.

583 :heavy_exclamation_mark: As `update` reads and rewrites POSEIDON.yml files, it may change their inner
584 order, layout or even content (e.g. if they have fields which are not in the [Poseidon package definition](#)). Create a
585 backup of the POSEIDON.yml file before running `update` if you are uncertain.

3 Inspection commands

3.1 List command

list lists packages, groups and individuals of the datasets you use, or of the packages available on the server.

Click here for command line details

```
Usage: trident list ((-d|--baseDir DIR) | --remote [--remoteURL ARG])
        (--packages | --groups | --individuals
        [-j|--jannoColumn JANNO_HEADER]) [--raw]
```

List packages, groups or individuals from local or remote Poseidon repositories

Available options:

```
-h,--help          Show this help text
-d,--baseDir DIR   a base directory to search for Poseidon Packages
                   (could be a Poseidon repository)
--remote           list packages from a remote server instead the local
                   file system
--remoteURL ARG    URL of the remote Poseidon server
                   (default: "https://c107-224.cloud.gwdg.de")
--packages         list all packages
--groups           list all groups, ignoring any group names after the
                   first as specified in the Janno-file
--individuals      list individuals
-j,--jannoColumn JANNO_HEADER
                   list additional fields from the janno files, using
                   the Janno column heading name, such as Country, Site,
                   Date_C14_Uncal_BP, Endogenous, ...
--raw             output table as tsv without header. Useful for piping
                   into grep or awk
--ignoreGeno       ignore SNP and GenoFile
```

To list packages from your local repositories, as seen above you can run

```
trident list -d ... -d ... --packages
```

This will yield a table like this

```
.------.------.------.
|           Title           |    Date    | Nr Individuals |
:=====:=====:=====:
| 2015_1000Genomes_1240K_haploid_pulldown | 2020-08-10 | 2535          |
| 2016_Mallick_SGDP1240K_diploid_pulldown | 2020-08-10 | 280           |
| 2018_BostonDatashare_modern_published   | 2020-08-10 | 2772          |
| ...                                     | ...        |               |
'------'------'-----'
```

so a nicely formatted table of all packages, their last update and the number of individuals in it.

627 To view packages on the remote server, instead of using directories to specify the locations of repositories on
628 your system, you can use `--remote` to show packages on the remote server. For example

629 `trident list --packages --remote`

630 will result in a view of all published packages in our [public online repository](#).

631 You can also list groups, as defined in the third column of EIGENSTRAT `.ind` files (or the first column of a
632 PLINK `.fam` file), and individuals with `--groups` and `--individuals` instead of `--packages`.

633 The `--individuals` flag provides a way to immediately access information from the `.janno`
634 files on the command line. This works with the `-j / --jannoColumn` option. For example adding
635 `--jannoColumn Country --jannoColumn Date_C14_Uncal_BP` to the commands above will add the `Country`
636 and the `Date_C14_Uncal_BP` columns to the respective output tables.

637 Note that if you want a less fancy table, for example because you want to load this into Excel, or pipe into
638 another command that cannot deal with the neat table layout, you can use the `--raw` option to output that
639 table as a simple tab-delimited stream.

640 3.2 Summarise command

641 `summarise` prints some general summary statistics for a given poseidon dataset taken from the `.janno` files.

642 [Click here for command line details](#)

643 Usage: `trident summarise (-d|--baseDir DIR) [--raw]`

644 Get an overview over the content of one or multiple Poseidon packages

645
646 Available options:

647 <code>-h,--help</code>	Show this help text
648 <code>-d,--baseDir DIR</code>	a base directory to search for Poseidon Packages (could be a Poseidon repository)
649 <code>--raw</code>	output table as tsv without header. Useful for piping 650 into <code>grep</code> or <code>awk</code>

652 You can run it with

653 `trident summarise -d ... -d ...`

654 which will show you context information like – among others – the number of individuals in the dataset, their
655 sex distribution, the mean age of the samples (for ancient data) or the mean coverage on the 1240K SNP array
656 in a table. `summarise` depends on complete `.janno` files and will silently ignore missing information for some
657 statistics.

658 You can use the `--raw` option to output the summary table in a simple, tab-delimited layout.

659 3.3 Survey command

660 `survey` tries to indicate package completeness (mostly focused on `.janno` files) for poseidon datasets.

661 [Click here for command line details](#)

662 Usage: `trident survey (-d|--baseDir DIR) [--raw]`

663 Survey the degree of context information completeness for Poseidon packages

664

665 Available options:

666	<code>-h,--help</code>	Show this help text
667	<code>-d,--baseDir DIR</code>	a base directory to search for Poseidon Packages (could be a Poseidon repository)
668		
669	<code>--raw</code>	output table as tsv without header. Useful for piping into grep or awk
670		

671 Running

672 `trident survey -d ... -d ...`

673 will yield a table with one row for each package. See `trident survey -h` for a legend which cell of this table
674 means what.

675 Again you can use the `--raw` option to output the survey table in a tab-delimited format.

676 **3.4 Validate command**

677 `validate` checks poseidon datasets for structural correctness.

678 Click here for command line details

679 Usage: `trident validate (-d|--baseDir DIR)`

680 Check one or multiple Poseidon packages for structural correctness

681

682 Available options:

683	<code>-h,--help</code>	Show this help text
684	<code>-d,--baseDir DIR</code>	a base directory to search for Poseidon Packages (could be a Poseidon repository)
685		
686	<code>--ignoreGeno</code>	ignore SNP and GenoFile
687	<code>--fullGeno</code>	test parsing of all SNPs (by default only the first 100 SNPs are probed)
688		
689	<code>--noExitCode</code>	do not produce an explicit exit code
690	<code>--ignoreDuplicates</code>	do not stop on duplicated individual names in the package collection
691		

692 You can run it with

693 `trident validate -d ... -d ...`

694 and it will either report a success (`Validation passed`) or failure with specific error messages to simplify
695 fixing the issues.

696 `validate` tries to ensure that each package in the dataset adheres to the [schema definition](#). Here is a list of
697 what is checked:

- 698 • Presence of the necessary files
- 699 • Full structural correctness of .bib and .janno file
- 700 • Superficial correctness of genotype data files by parsing the first 100 SNPs. A full check that parses all
701 SNPs can be run with the `--fullGeno` option
- 702 • Correspondence of BibTeX keys in .bib and .janno

703 • Correspondence of individual and group IDs in .janno and genotype data files

704 In fact much of this validation already runs as part of the general package reading pipeline invoked for many
705 trident subcommands (e.g. `forge`). `validate` is meant to be more thorough, though, and will explicitly fail if
706 even a single package is broken.

707 Remember to run it with `--logMode VerboseLog` to get more information if the output is not sufficient to
708 debug an issue.